

Instruction Set Architectures and Closures

CIS400 (Compiler Construction)

Kris Micinski, Fall 2021

- In the last lecture, we talked about closure conversion
- Today, I want to continue to motivate why closure conversion is a key step in compilation of functional languages

Key ideas:

- ➊ Closure conversion **eliminates lambdas** by turning them into “C-style” functions
 - ➌ **Why?**
 - ➌ Computers really good at calling C-style functions!
 - ➌ Processor has a native **call** instruction to invoke first-order procedures (C functions)
- ➋ Closure invocation then becomes *dynamic dispatch*

- In the last lecture, we talked about closure conversion
- Today, I want to continue to motivate why closure conversion is a key step in compilation of functional languages

Key ideas:

In today's lecture, I hope to motivate this...

- Closure conversion **eliminates lambdas** by turning them into “C-style” functions
 - **Why?**
 - Computers really good at calling C-style functions!
 - Processor has a native **call** instruction to invoke first-order procedures (C functions)
- Closure invocation then becomes *dynamic dispatch*



Course Detour — Assembly Crash Course

By which I mean x86-64 assembly...

Note to readers...

In this set of slides, I give a brief detour of how x86-64 assembly works. In this course we will not use x86-64 assembly, we will use LLVM assembly. LLVM is a higher-level assembly than x86-64, and while many of the concepts will be similar, they will not be identical.

ISA



- The **I**nstruction **S**et **A**rchitecture is the term used to refer to the native language executed by a (hardware) processor
 - vs. (say) a *bytecode* executed by a virtual machine (such as the JVM)
- All processors execute a set of very basic commands called *instructions*
- These instructions are kind of like ANF/CPS—they atomically execute one instruction at a time

Mov, sub, etc... are destination-last here, but there is no standard and in practice you should get used to figuring out if the instruction is destination-first or destination-last...

Opcode — instruction name

```
pushl %ebp ; save calling function's stack frame (ebp)
movl %esp, %ebp ; make a new stack frame on top of our caller's stack
subl $4, %esp ; allocate 4 bytes of stack space for local variables
```

Operands—atomic arguments to instructions

Operands must be atomic

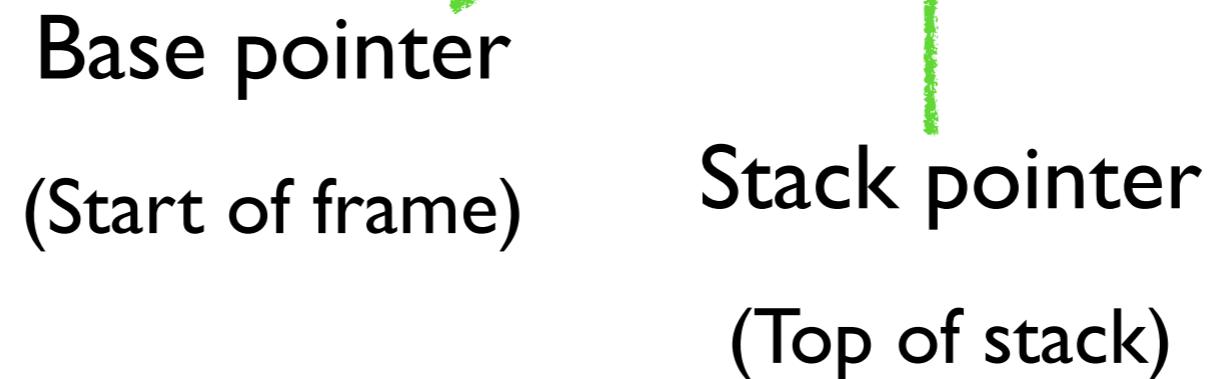
- There are several kinds of atomic values in assembly:
 - Literal values
 - Registers (high-speed temporary variables close to process)
 - RAM (may be backed by cache) such as the stack

Registers

Originally, 8-bit registers: al, bl, cl, dl

Traditionally, x86 architectures only had **four** 16-bit general purpose registers: ax, bx, cx, dx

Also other registers: bp, sp, di, si

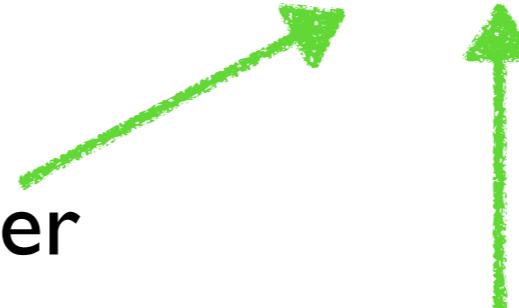


Originally, 8-bit registers: al, bl, cl, dl

Traditionally, x86 architectures only had **four** 16-bit general purpose registers: ax, bx, cx, dx

Also other registers: bp, sp, di, si

Base pointer
(Start of frame)



Stack pointer
(Top of stack)

IP: instruction pointer

Points at current instruction,
incremented after each instruction

FLAGS: holds flags

Set on subtraction, comparison, etc..

Traditionally, x86 architectures only had **four** 16-bit general purpose registers: ax, bx, cx, dx

Also other registers: bp, sp, di, si

As time progressed, also added 32-bit registers:
eax, ebx, ecx, edx

In past few years, 64-bit registers: rax, rbx, rcx, rdx
(Also 64-bit versions: rip, etc..)

We'll pretty much exclusively use
64-bit registers!

Note RAX is an **extension** of EAX

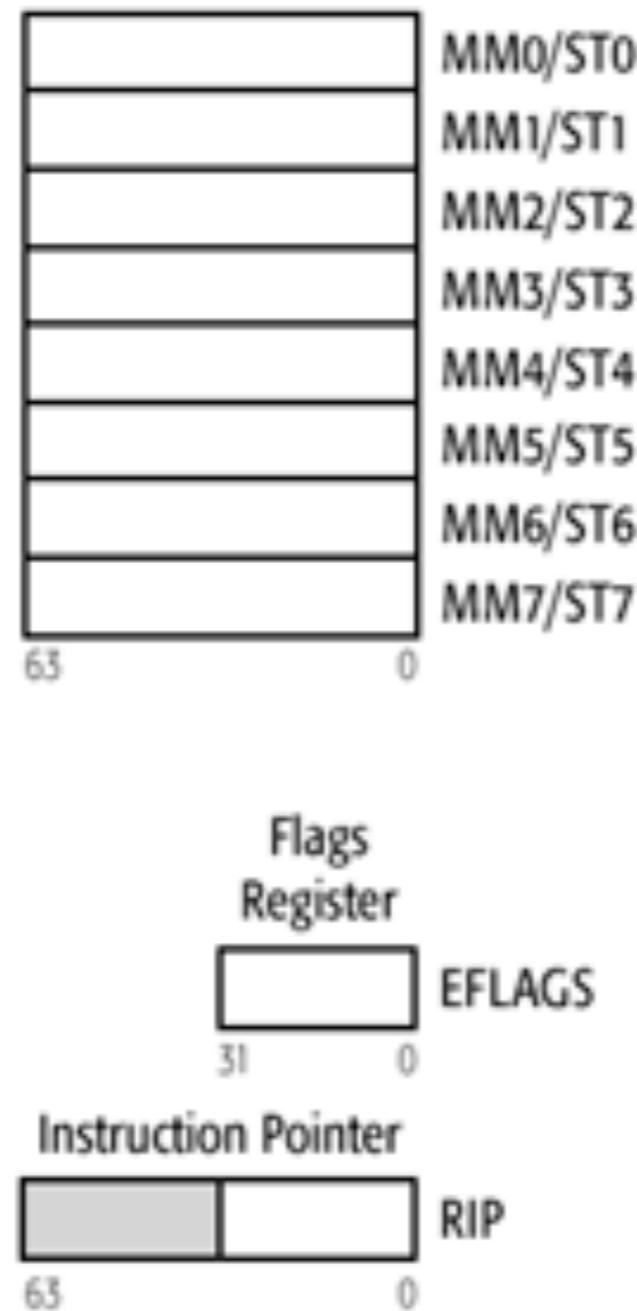


If you change EAX, you change lower 32 bits of RAX

General-Purpose Registers (GPRs)

	RAX
	RBX
	RCX
	RDX
	RBP
	RSI
	RDI
	RSP
63	R8
	R9
	R10
	R11
	R12
	R13
	R14
	R15

Multimedia Extension and Floating-Point Registers



Streaming SIMD Extension (SSE) Registers

XMM0
XMM1
XMM2
XMM3
XMM4
XMM5
XMM6
XMM7
XMM8
XMM9
XMM10
XMM11
XMM12
XMM13
XMM14
XMM15



Legacy x86 Registers, supported in all modes



Register Extensions, supported in 64-Bit Mode

General-Purpose Registers (GPRs)

	RAX
	RBX
	RCX
	RDX
	RBP
	RSI
	RDI
	RSP
63	R8
	R9
	R10
	R11
	R12
	R13
	R14
63	R15
	0

Multimedia Extension and Floating-Point Registers

MM0/ST0
MM1/ST1
MM2/ST2
MM3/ST3
MM4/ST4
MM5/ST5
MM6/ST6
MM7/ST7

63 0

Flags
Register

EFLAGS

31 0

Instruction Pointer

63 0

RIP

Streaming SIMD Extension (SSE) Registers

XMM0
XMM1
XMM2
XMM3
XMM4
XMM5
XMM6
XMM7
XMM8
XMM9
XMM10
XMM11
XMM12
XMM13
XMM14
XMM15

0

127



Legacy x86 Registers, supported in all modes

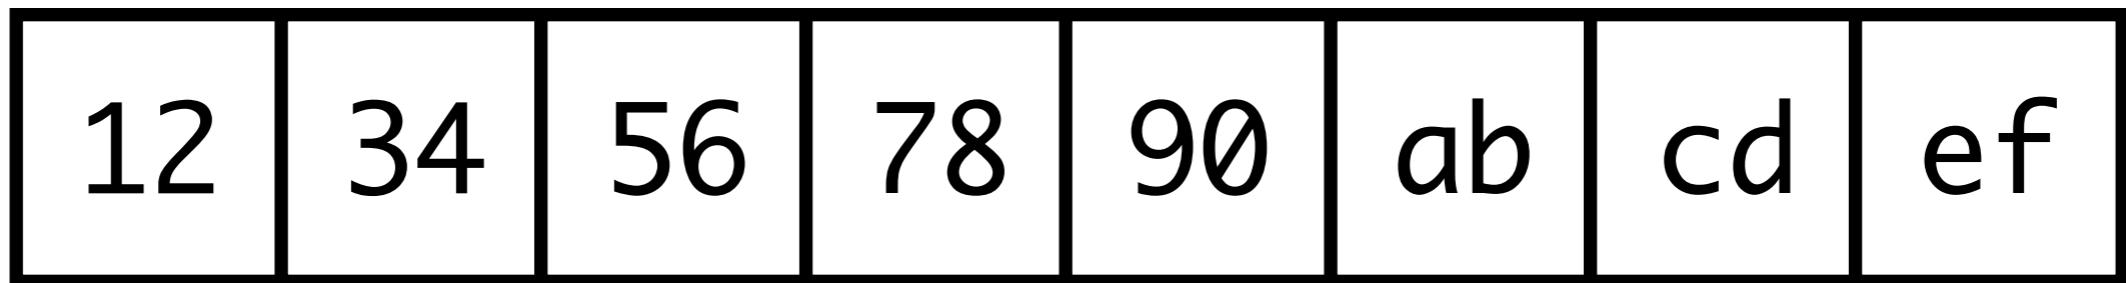


Register Extensions, supported in 64-Bit Mode

Special regs: floating-point / matrix ops

507001.eps

To represent `0x1234567890abcdef`



Most Significant Byte

Least Significant Byte

x86 is a **little-endian** architecture

If an n-byte value is stored at addresses a to $a+(n-1)$ in memory,
byte a will hold the **least significant byte**

0x1234567890abcdef

Exercise with partner

Instructions

Binary code is made up of giant sequences of “instructions”

Modern Intel / AMD chip has hundreds of them, some very complex

Moving memory around

Arithmetic

Branch / If

Matrix operations

Atomic-Instructions

Transactional memory instructions

Encoded as binary (as you may have seen from
hardware-design course)

We (humans) write in a format named “assembly”

Confusingly: two types of assembly

AT&T

movq 5, %rax

Intel

mov rax, 5

Several addressing modes

“Move the value from register `rax` into the register `rbx`”

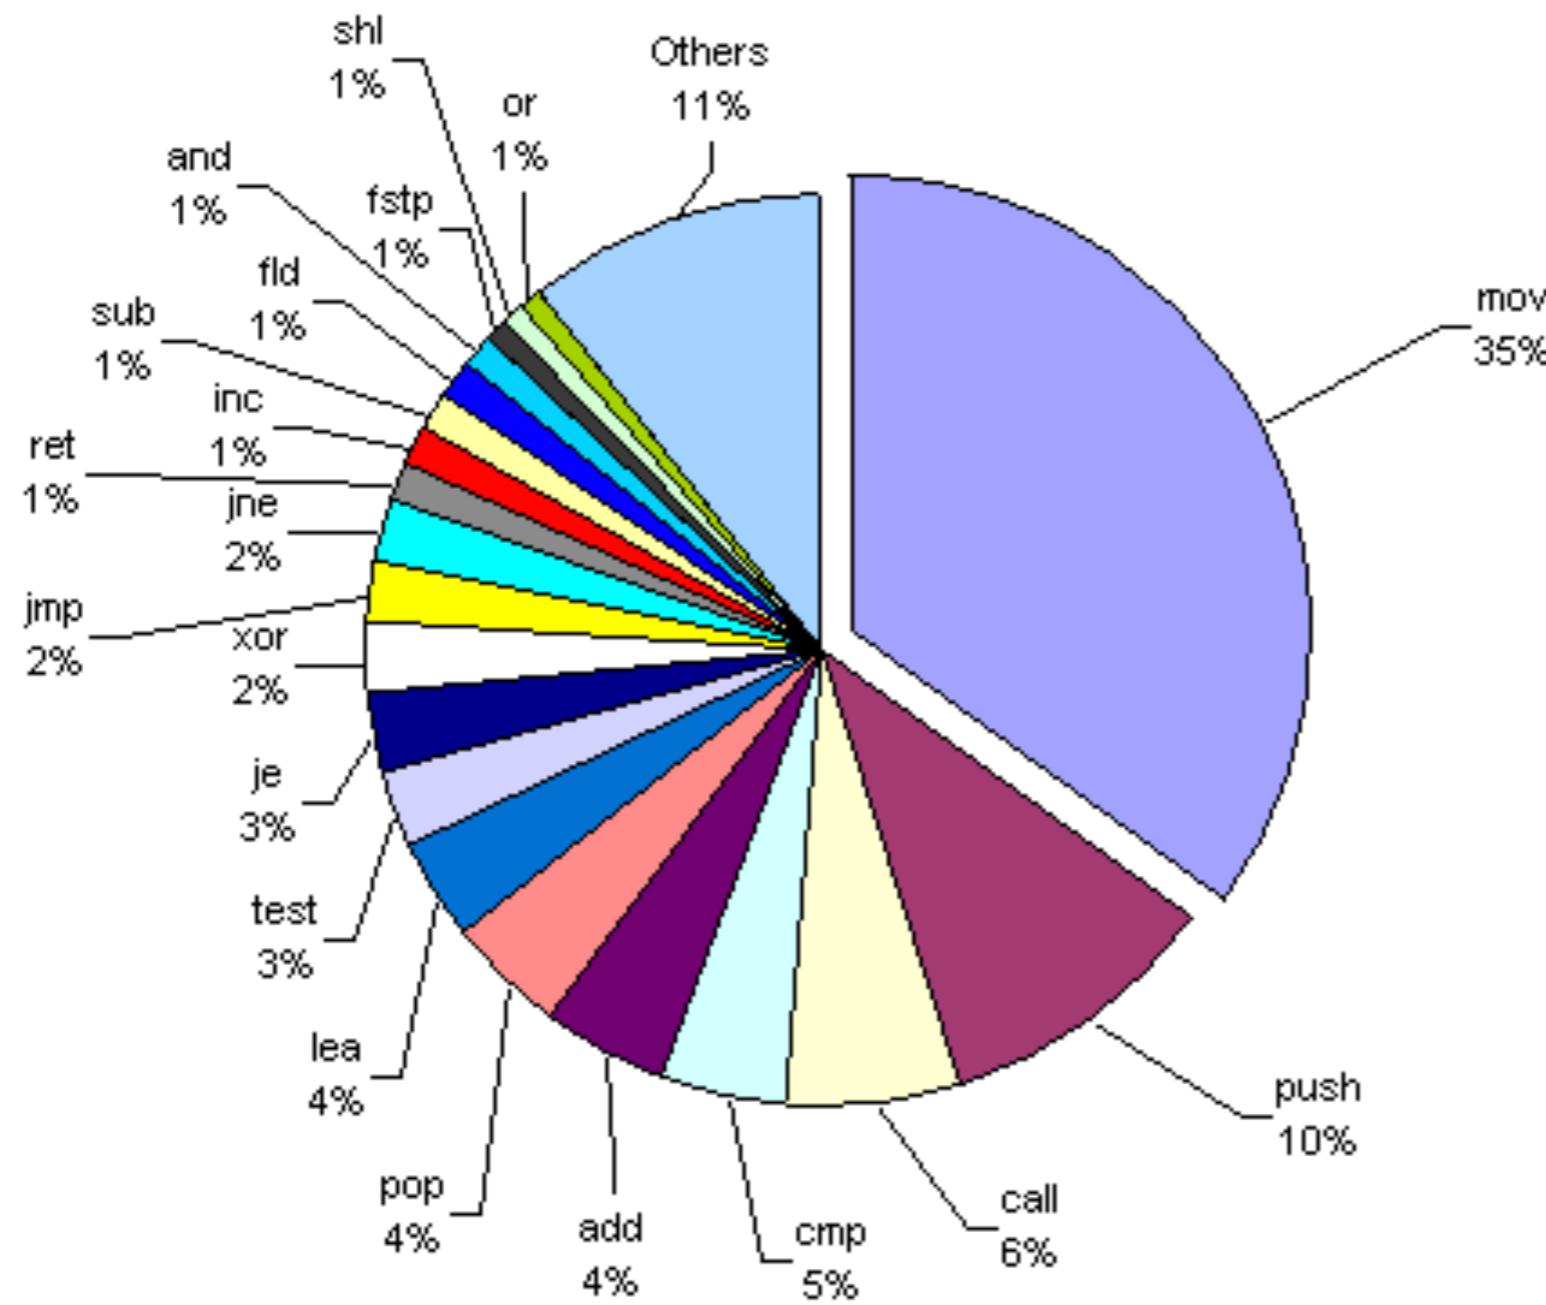
Opcode name

`mov %rax, %rbx`

Destination

Source

Top 20 instructions of x86 architecture



Plurality of instructions
are **movs**

Then **push**

Then **call**

Memory: a **giant chunk of bytes**

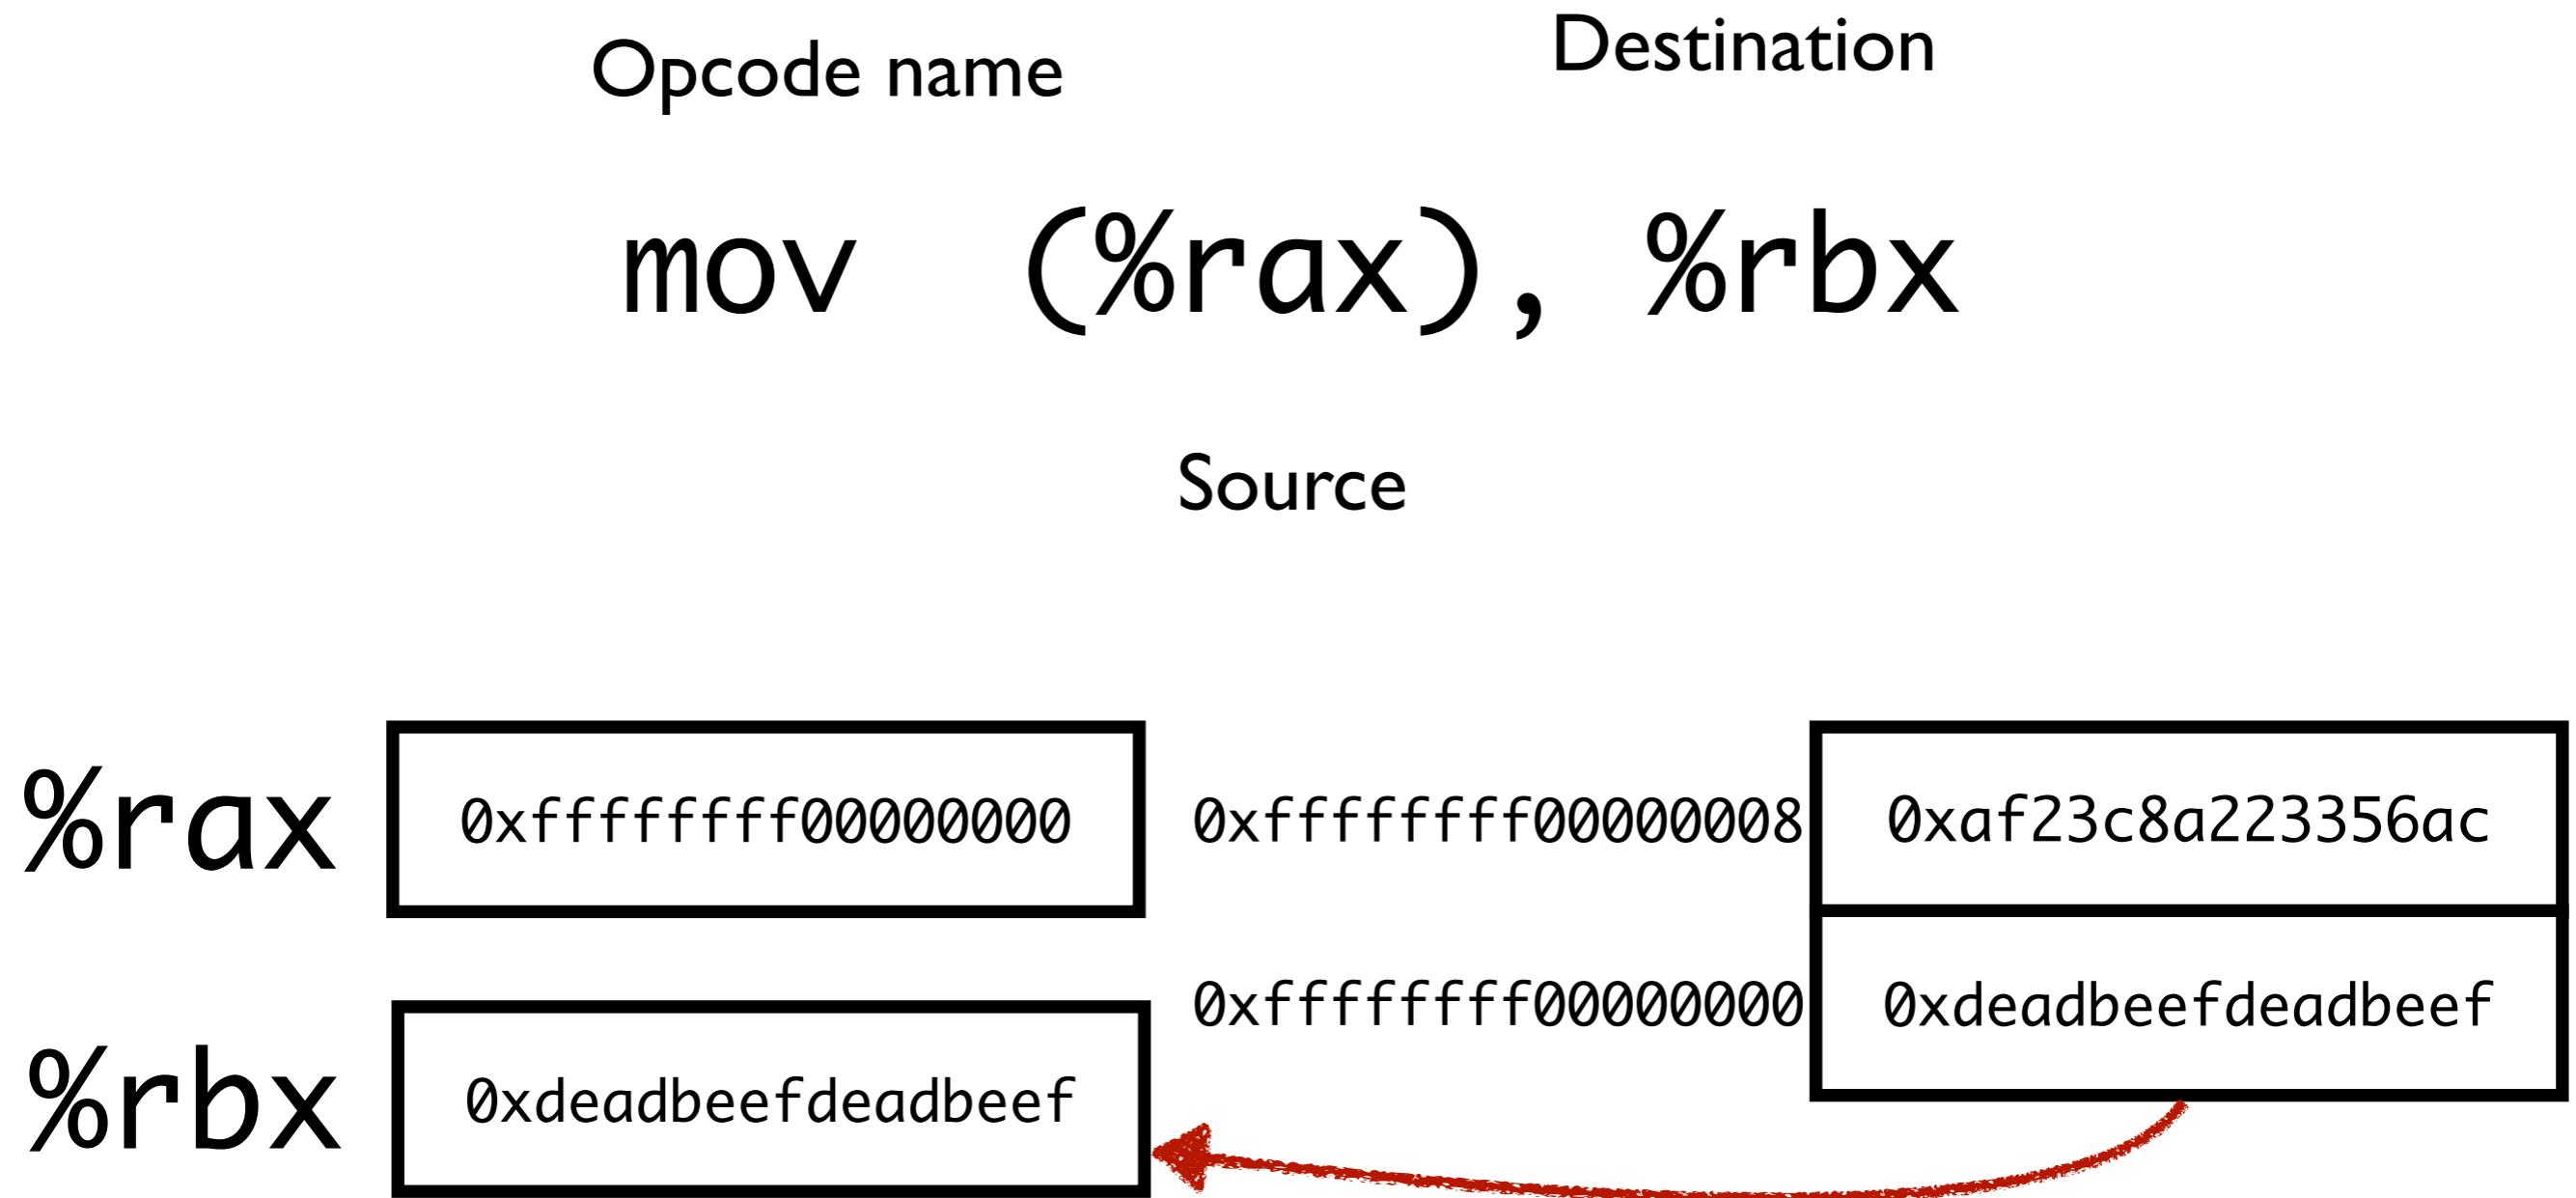
You can read from it and write to it in 1/2/4/8/16-byte increments

```
mov    (%rax), %rbx
```

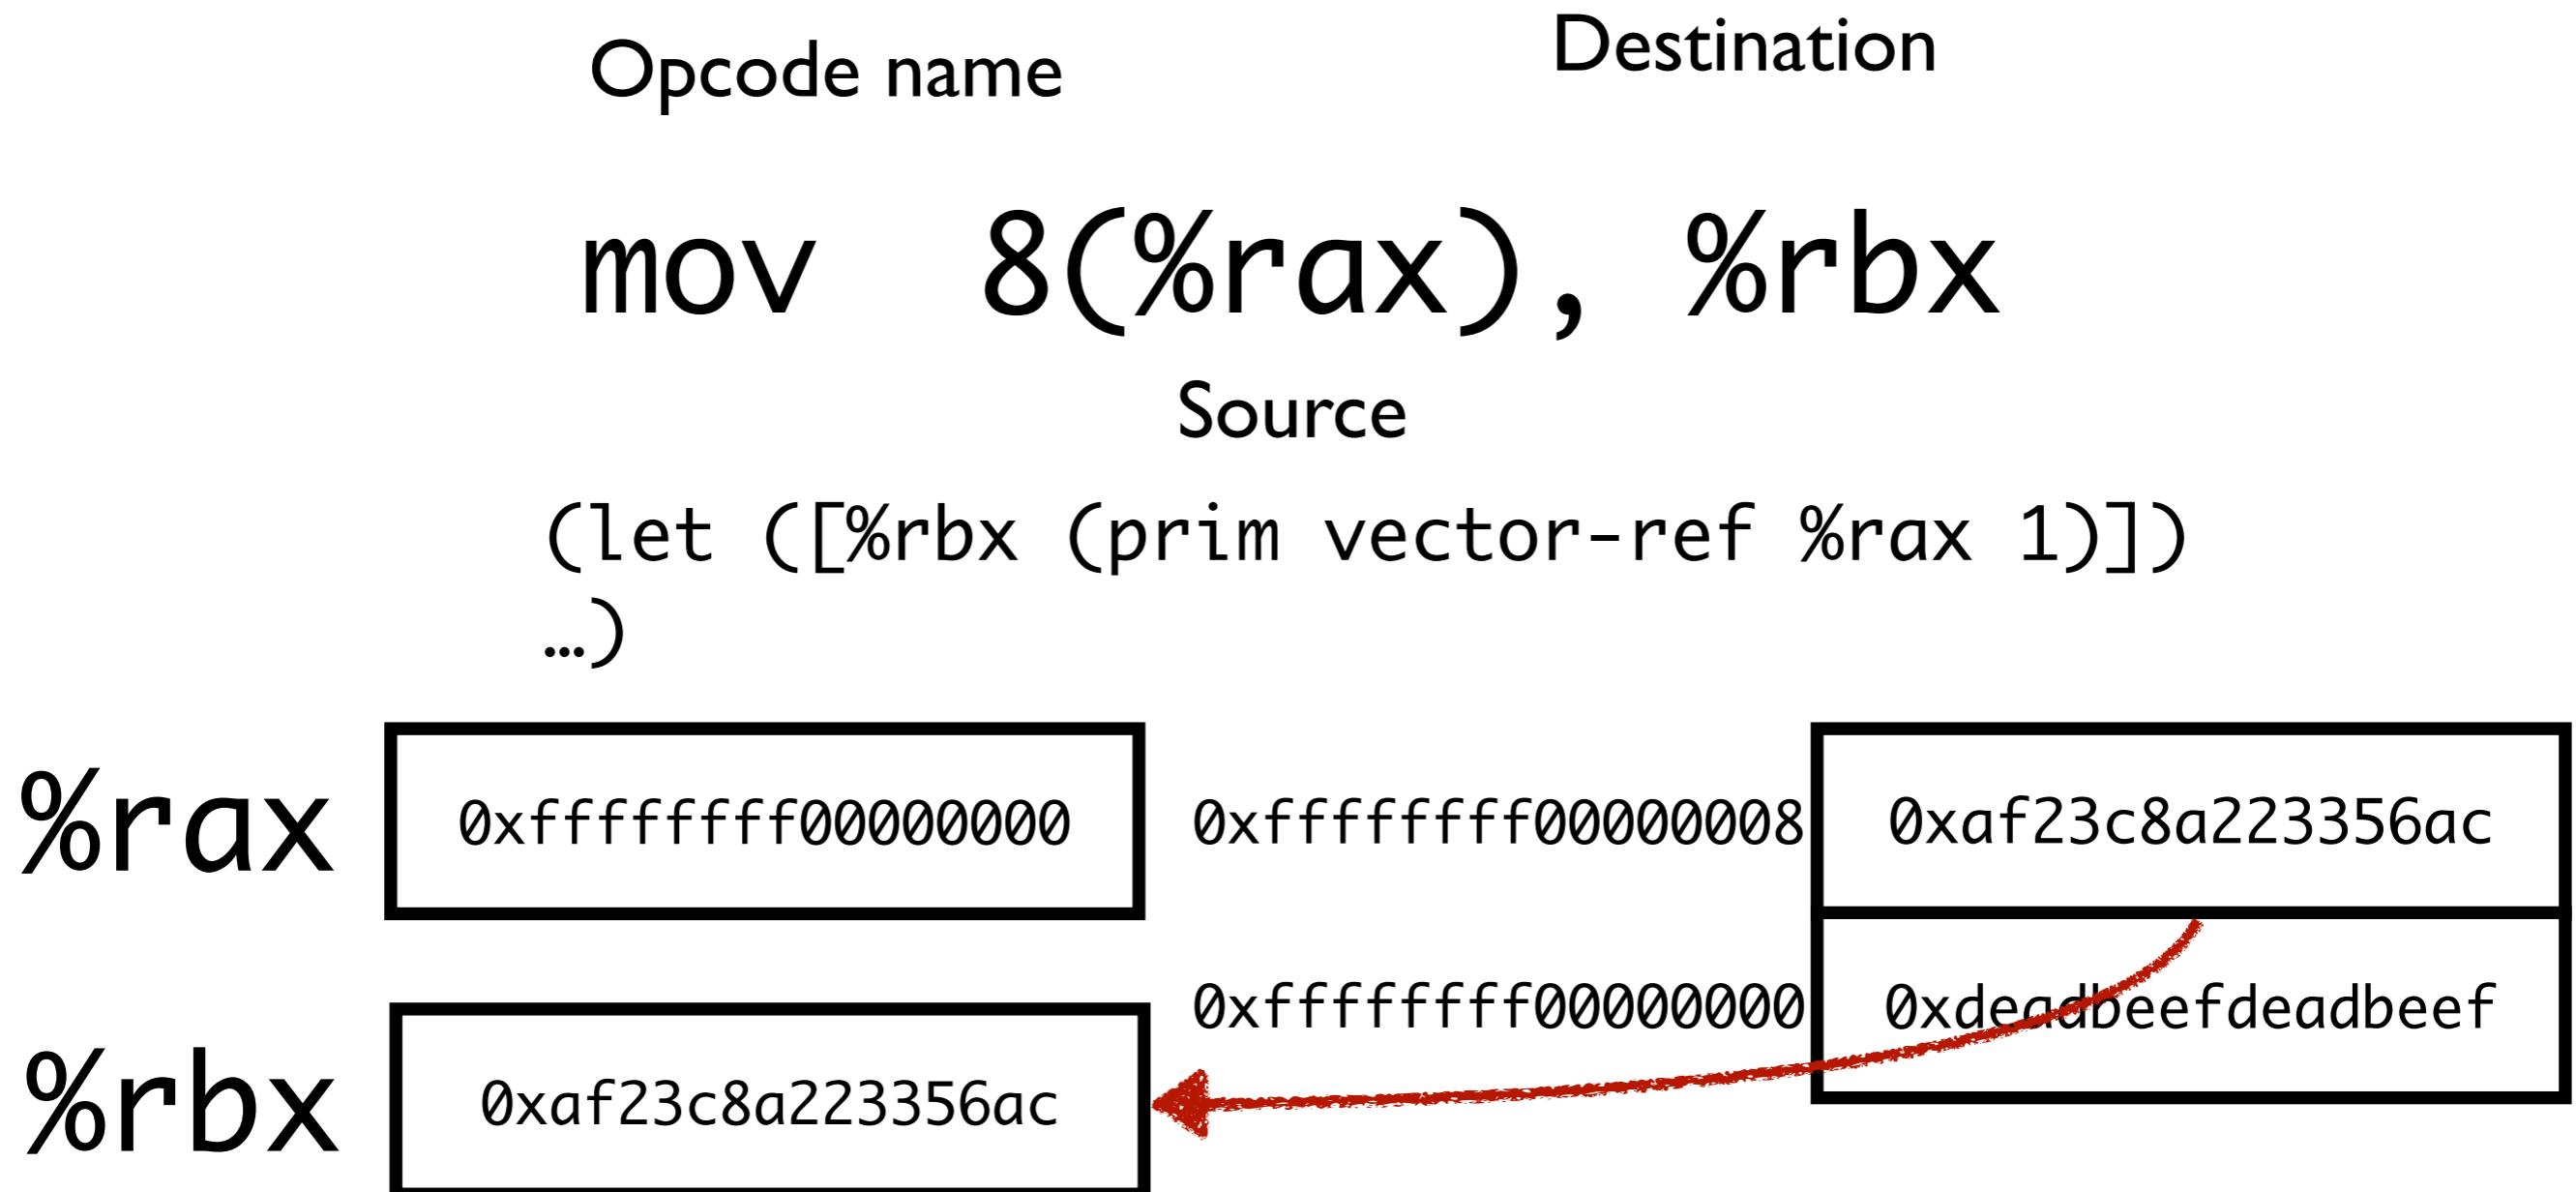
“Move the value **at address %rax** into register **%rbx**”

Opcode name	Destination	Source
mov (%rax), %rbx		
%rax	0xffffffff00000008	0xaf23c8a223356ac
%rbx	0xffffffff00000000	0xdeadbeefdeadbeef

“Move the value **at address %rax** into register **%rbx**”



“Move the value at address %rax+8 into register %rbx”



A few other more complicated ones that allow you to add registers, offsets, etc...

Different instructions allow different addressing-modes

Caching

- Accessing RAM is slow, much slower than registers (~1-2 orders of magnitude, depending on how recently used)
- All memory access is backed by a **cache**, which mirrors recently-accessed pieces of RAM in an on-chip cache

Caching (contd...)

- When *cache miss* occurs (data not in cache), processor pulls the data's enclosing **cache line** (64-byte-aligned sequence of memory around the data) into the cache
 - Thus, accessing data within that cache line will be *fast*
- Designing for cache friendliness can be crucial
 - One **take-away** point for **all** students—if you touch data once, all data sharing the same cache line will be extremely fast to use
 - ***One course lesson — Situate data that will be used together on the same cache line!***

RISC vs. CISC

- x86_64 (i.e., Intel/AMD-style assembly) is a **Complex Instruction Set Computer**—a huge mess of complex instructions to do tons of random things
 - This design is crummy for lots of reasons, and modern machines decode x86_64 into an internal RISC anyway, wasting lot of silicon area
- ARM is a competitor to x86_64, and it is a **Reduced Instruction Set Computer (RISC)** instruction set
 - ARM has taken over the world of micro devices, phones, tablets, and now laptops and server chips
 - Much lower power, simpler instruction set, conventional wisdom says it will be less fast (will it? Likely not, server-class ARM chips now power world's fastest supercomputer by Fujitsu)

Example ARM code...

```
int total;  
int i;  
  
total = 0;  
for (i = 10; i > 0; i--) {  
    total += i;  
}
```

```
        MOV R0, #0          ; R0 accumulates total  
        MOV R1, #10         ; R1 counts from 10 down to 1  
again   ADD R0, R0, R1  
        SUBS R1, R1, #1  
        BNE again  
halt    B halt          ; infinite loop to stop computation
```

Memory is divided into different regions

C stack (grows up and down)

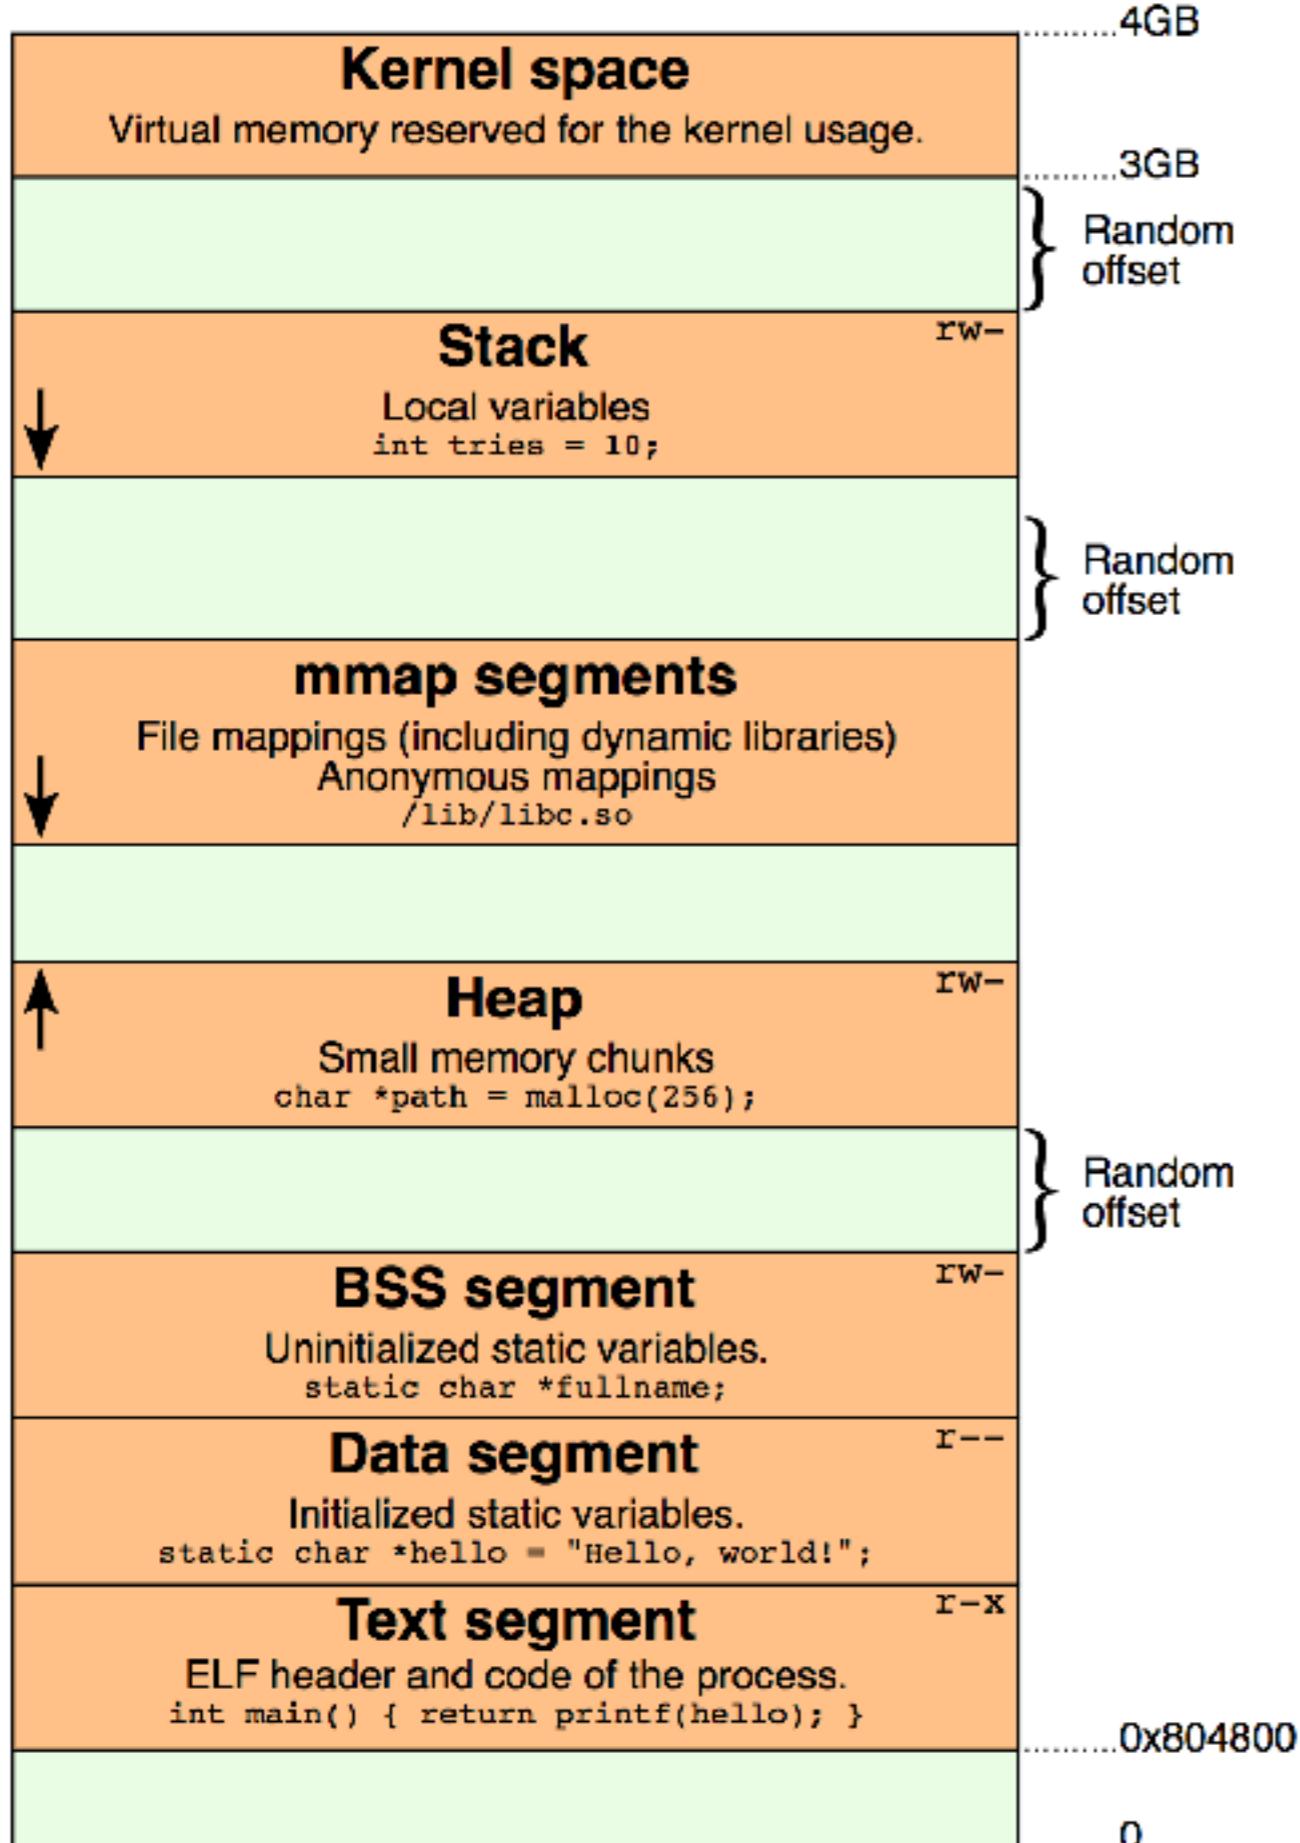
Heap (new/malloc'd data)

Static variables

Code (text)

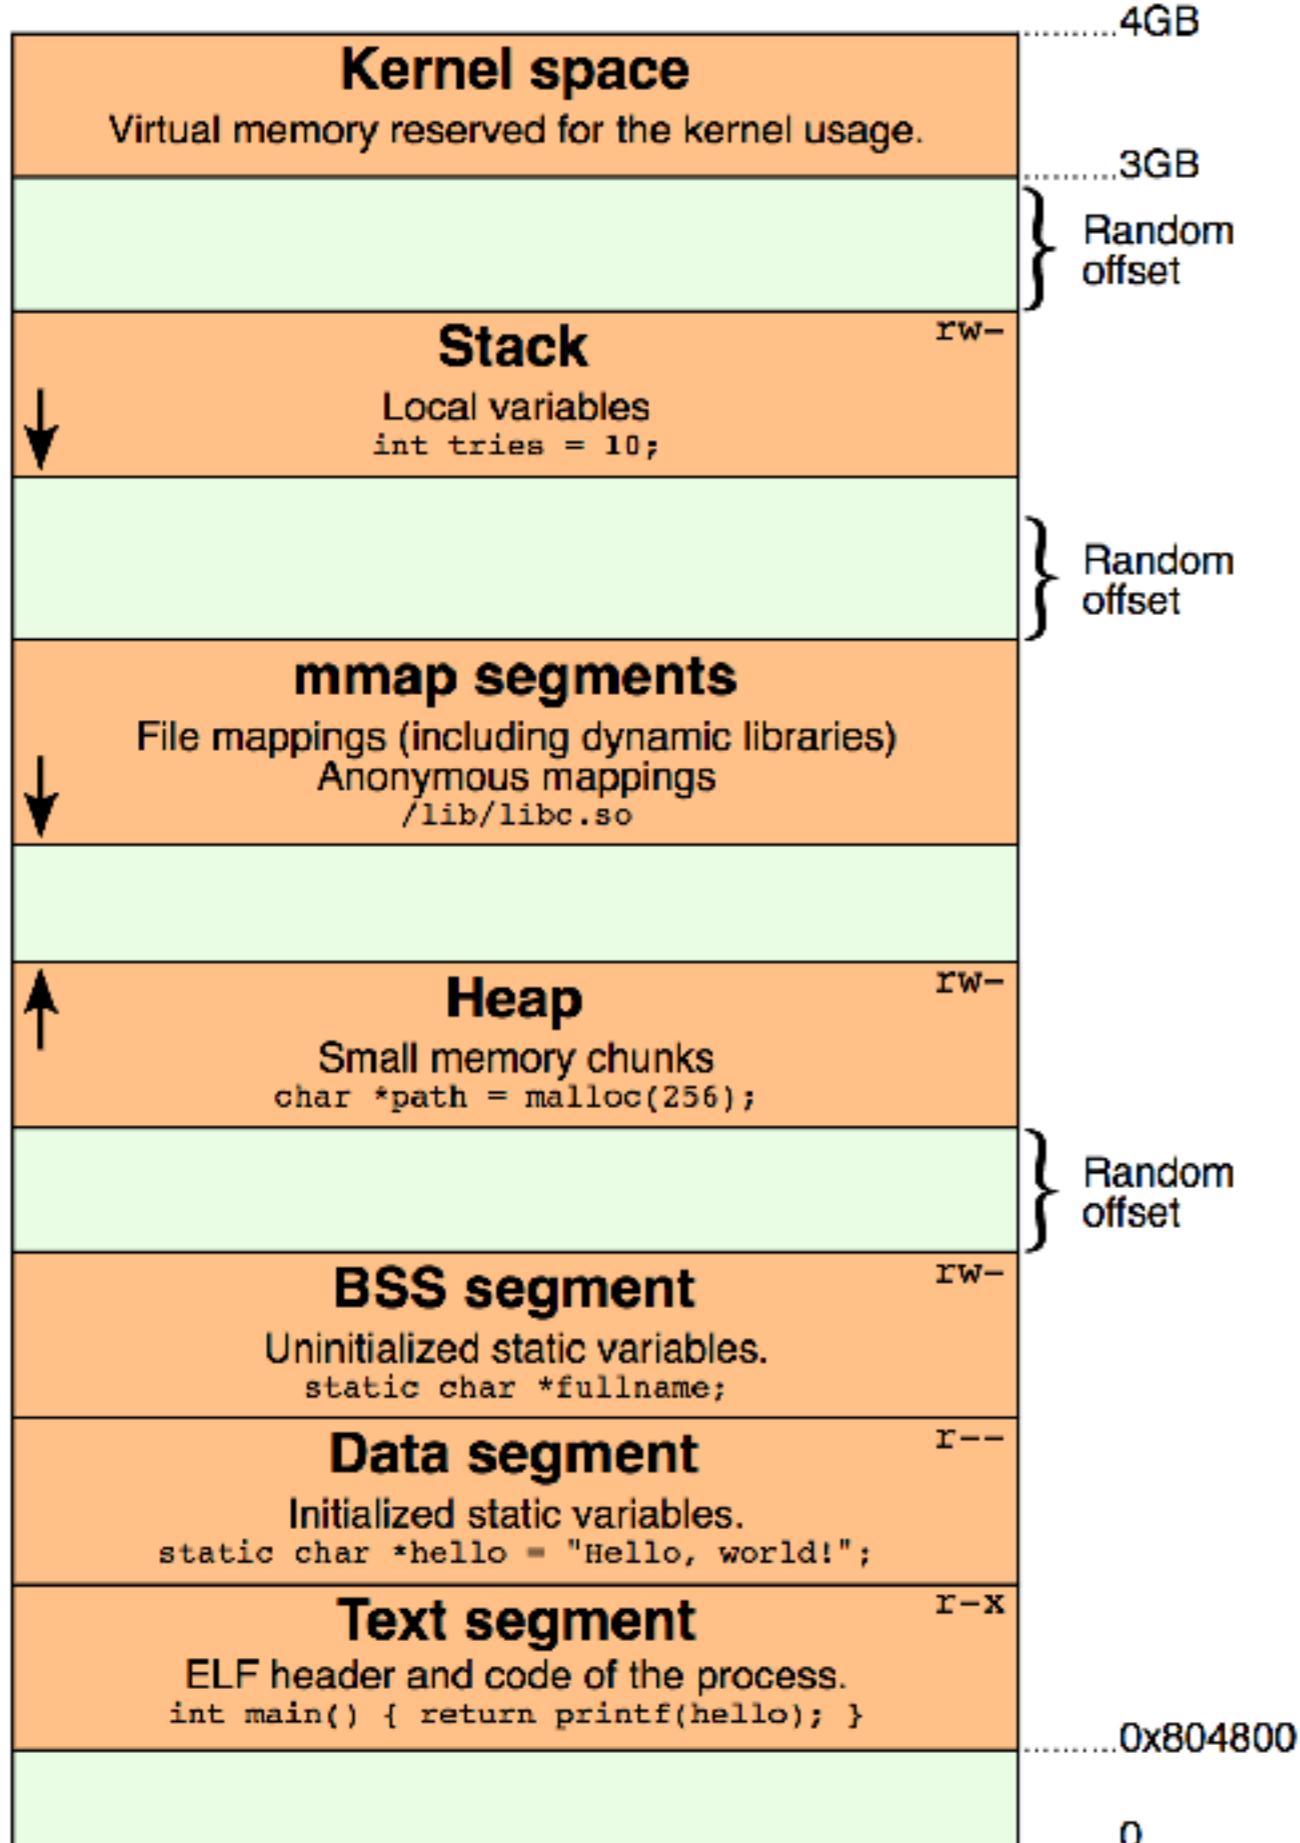
...

OS loads these into different **segments** for a variety of (mostly security-related) reasons



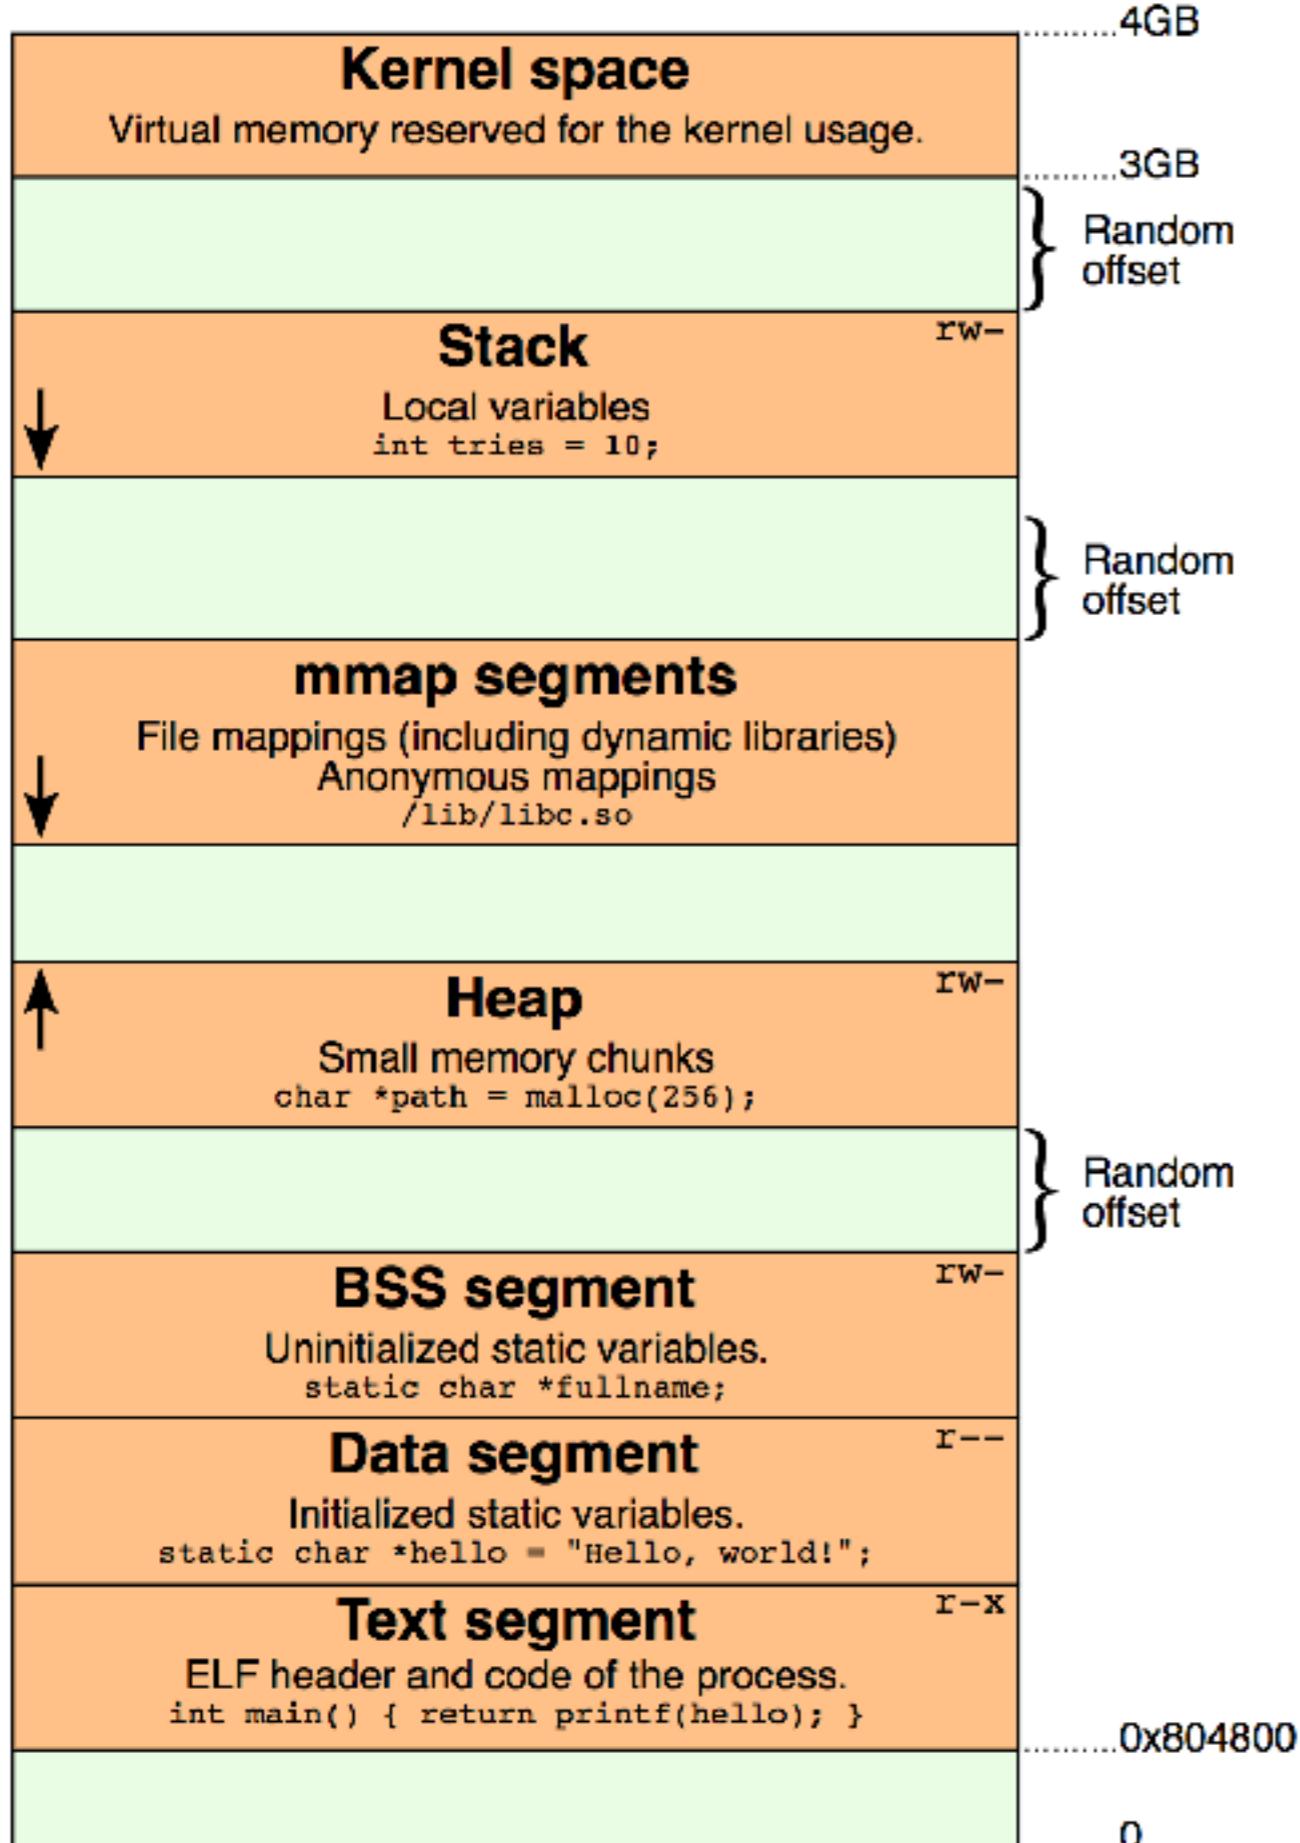
Kernel memory

Your OS uses it



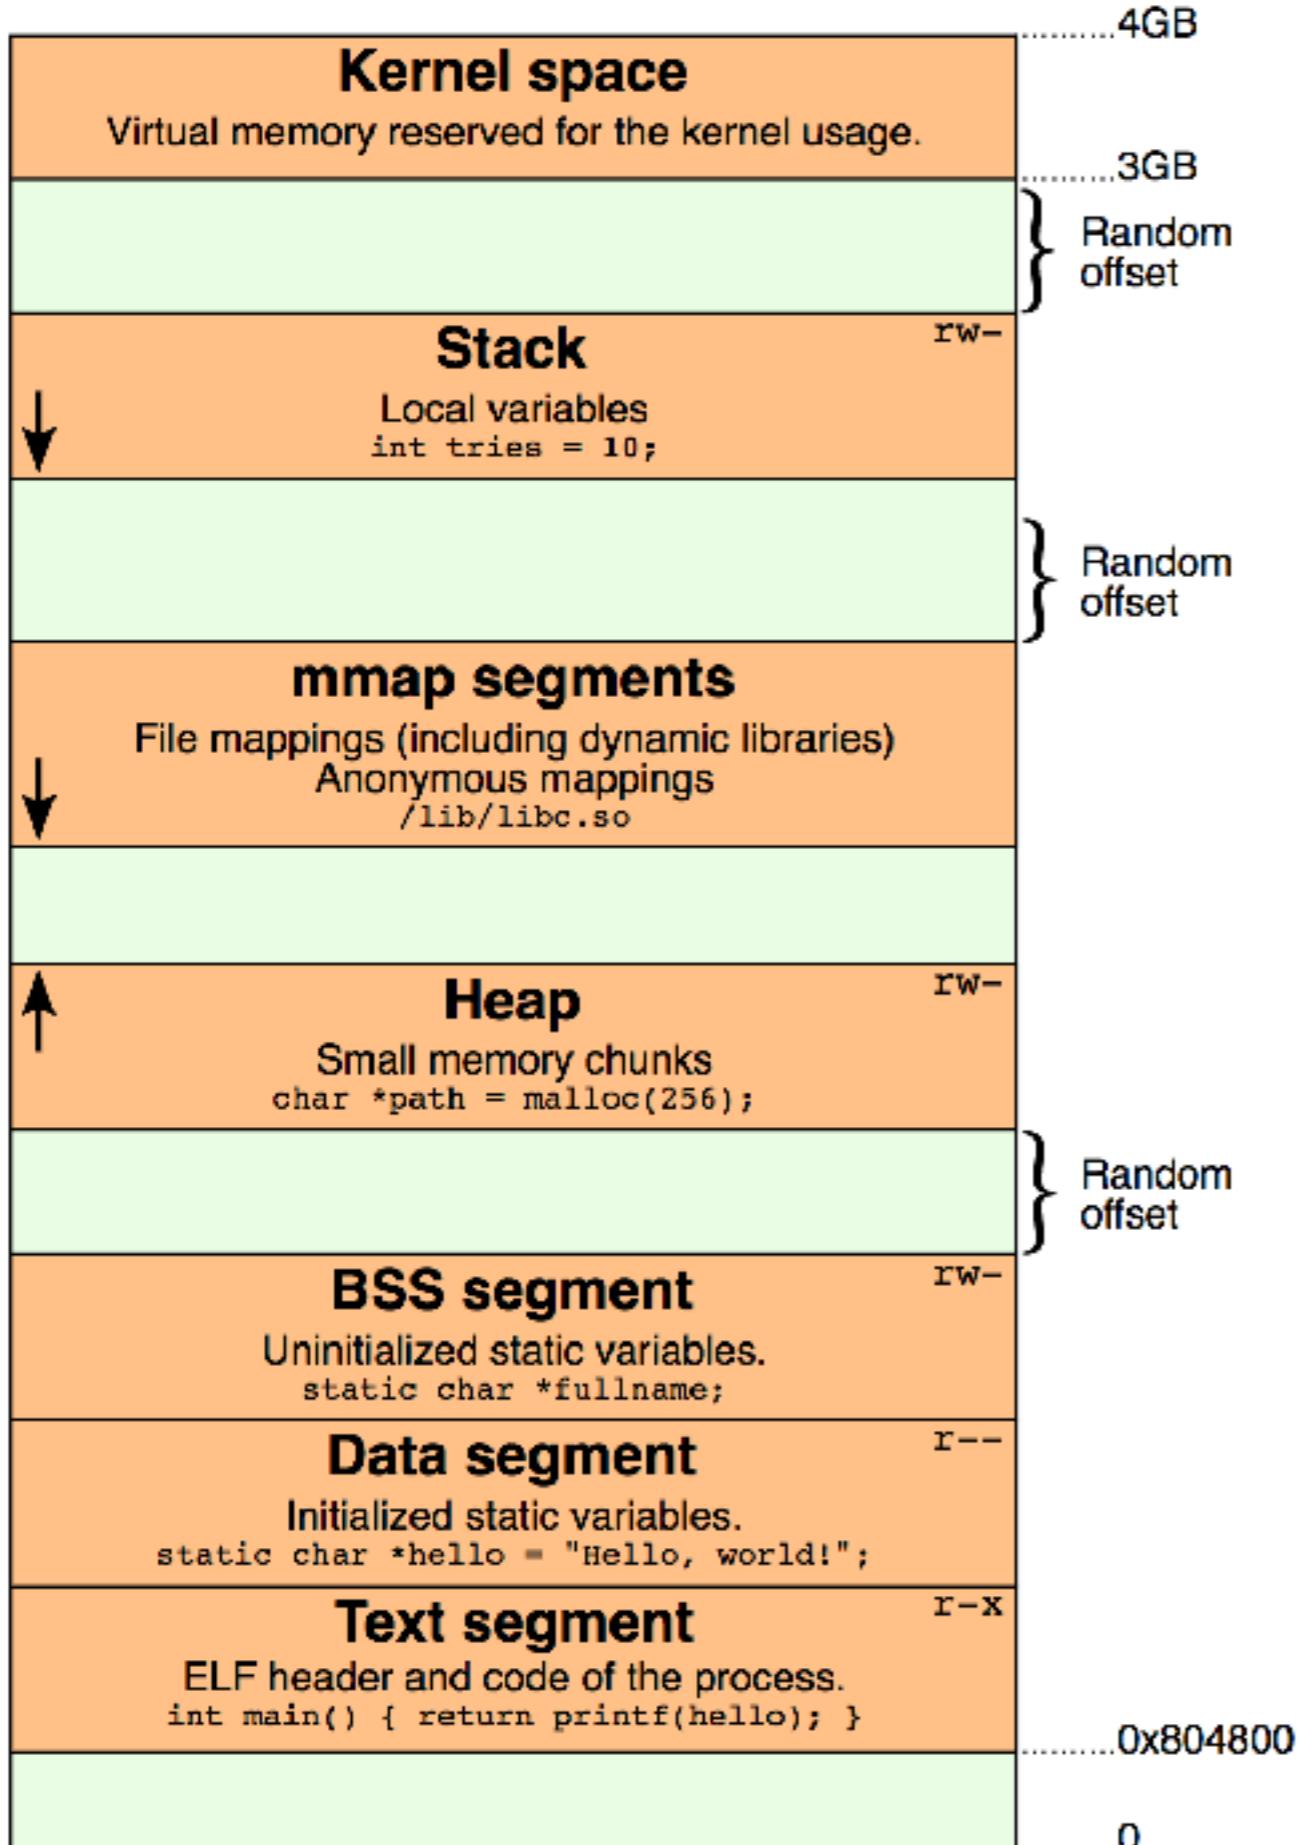
Stack: push / pop

Very important:
The stack grows **down**



mmap segments

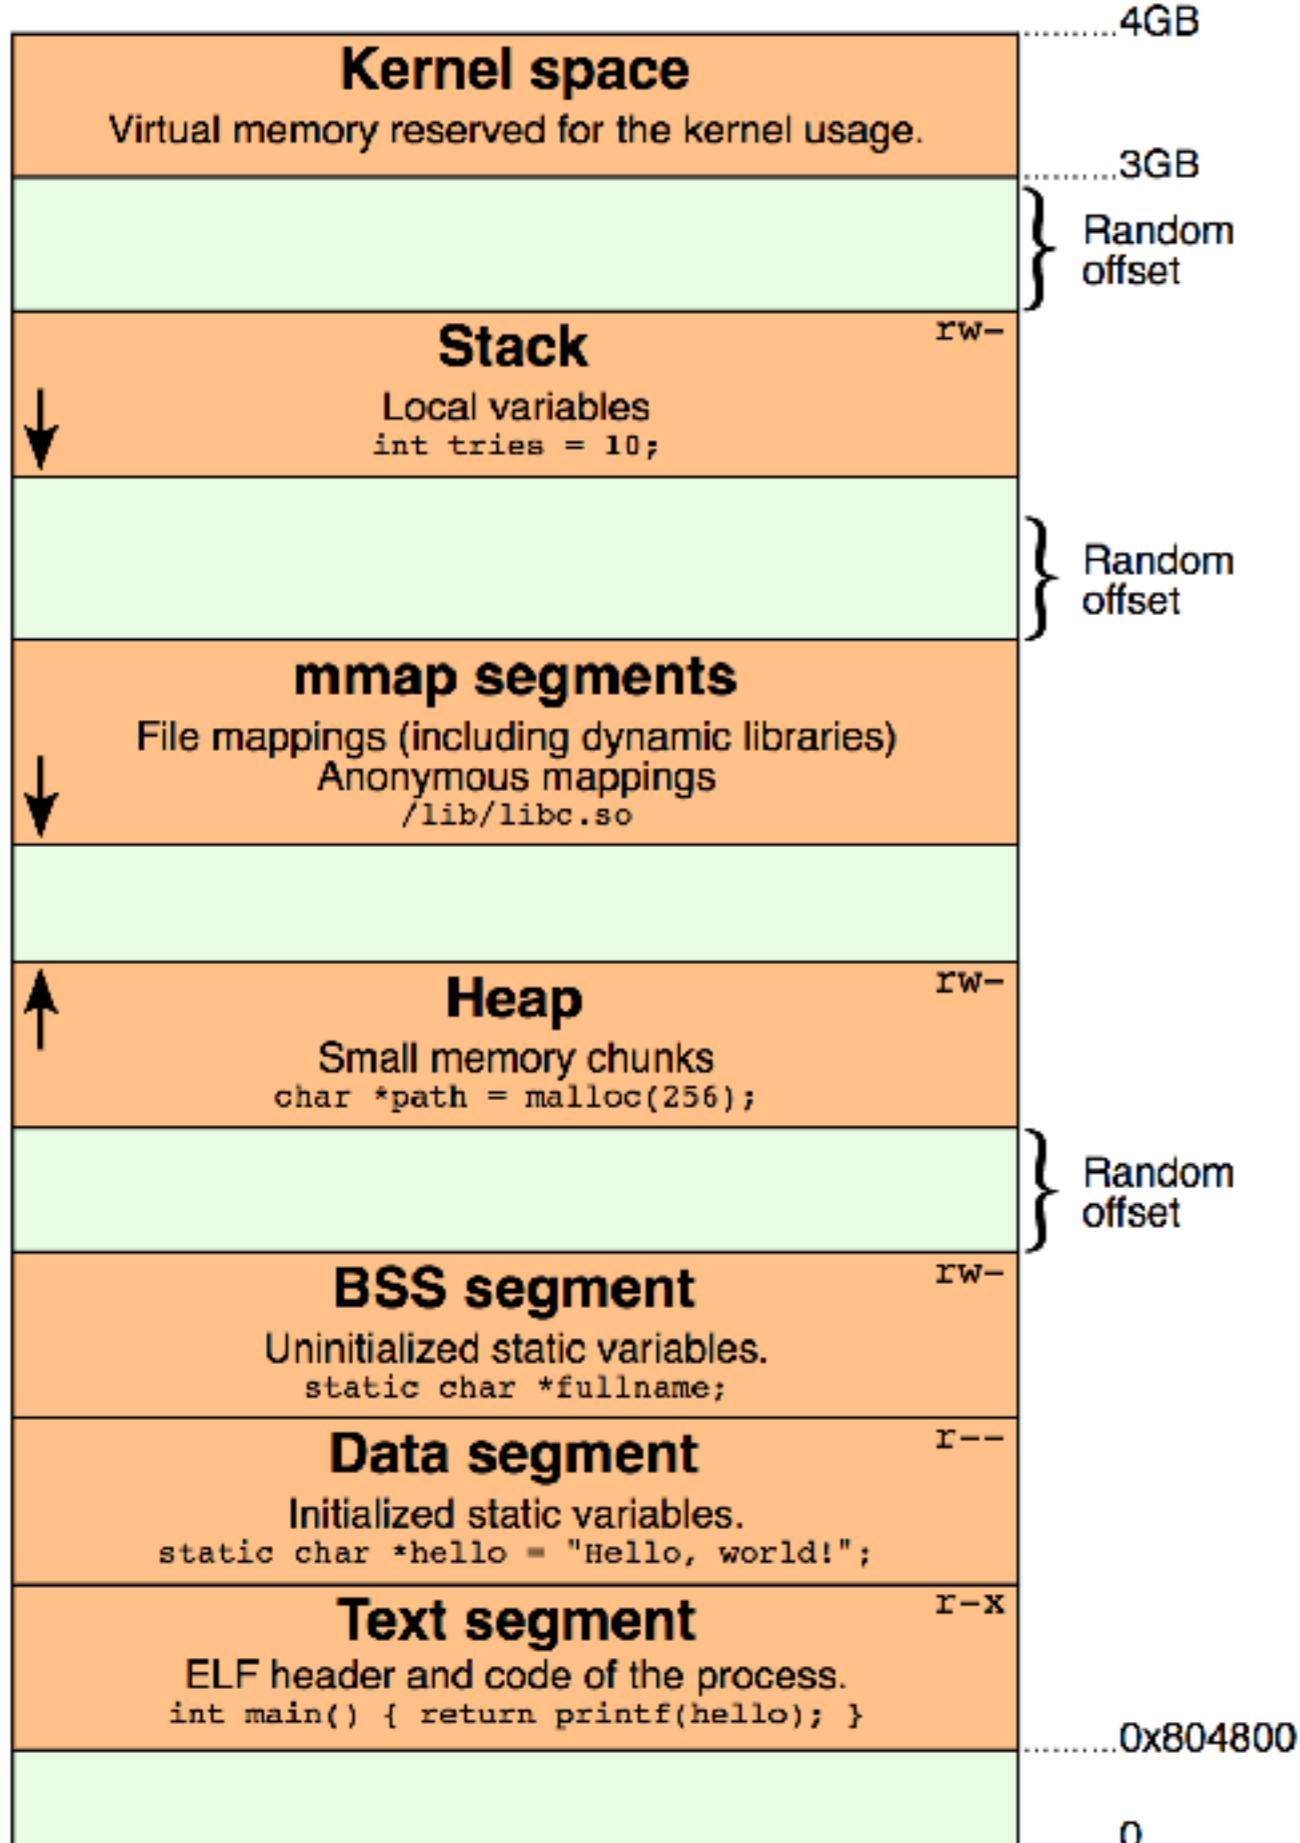
Allows you to **map** a file to memory



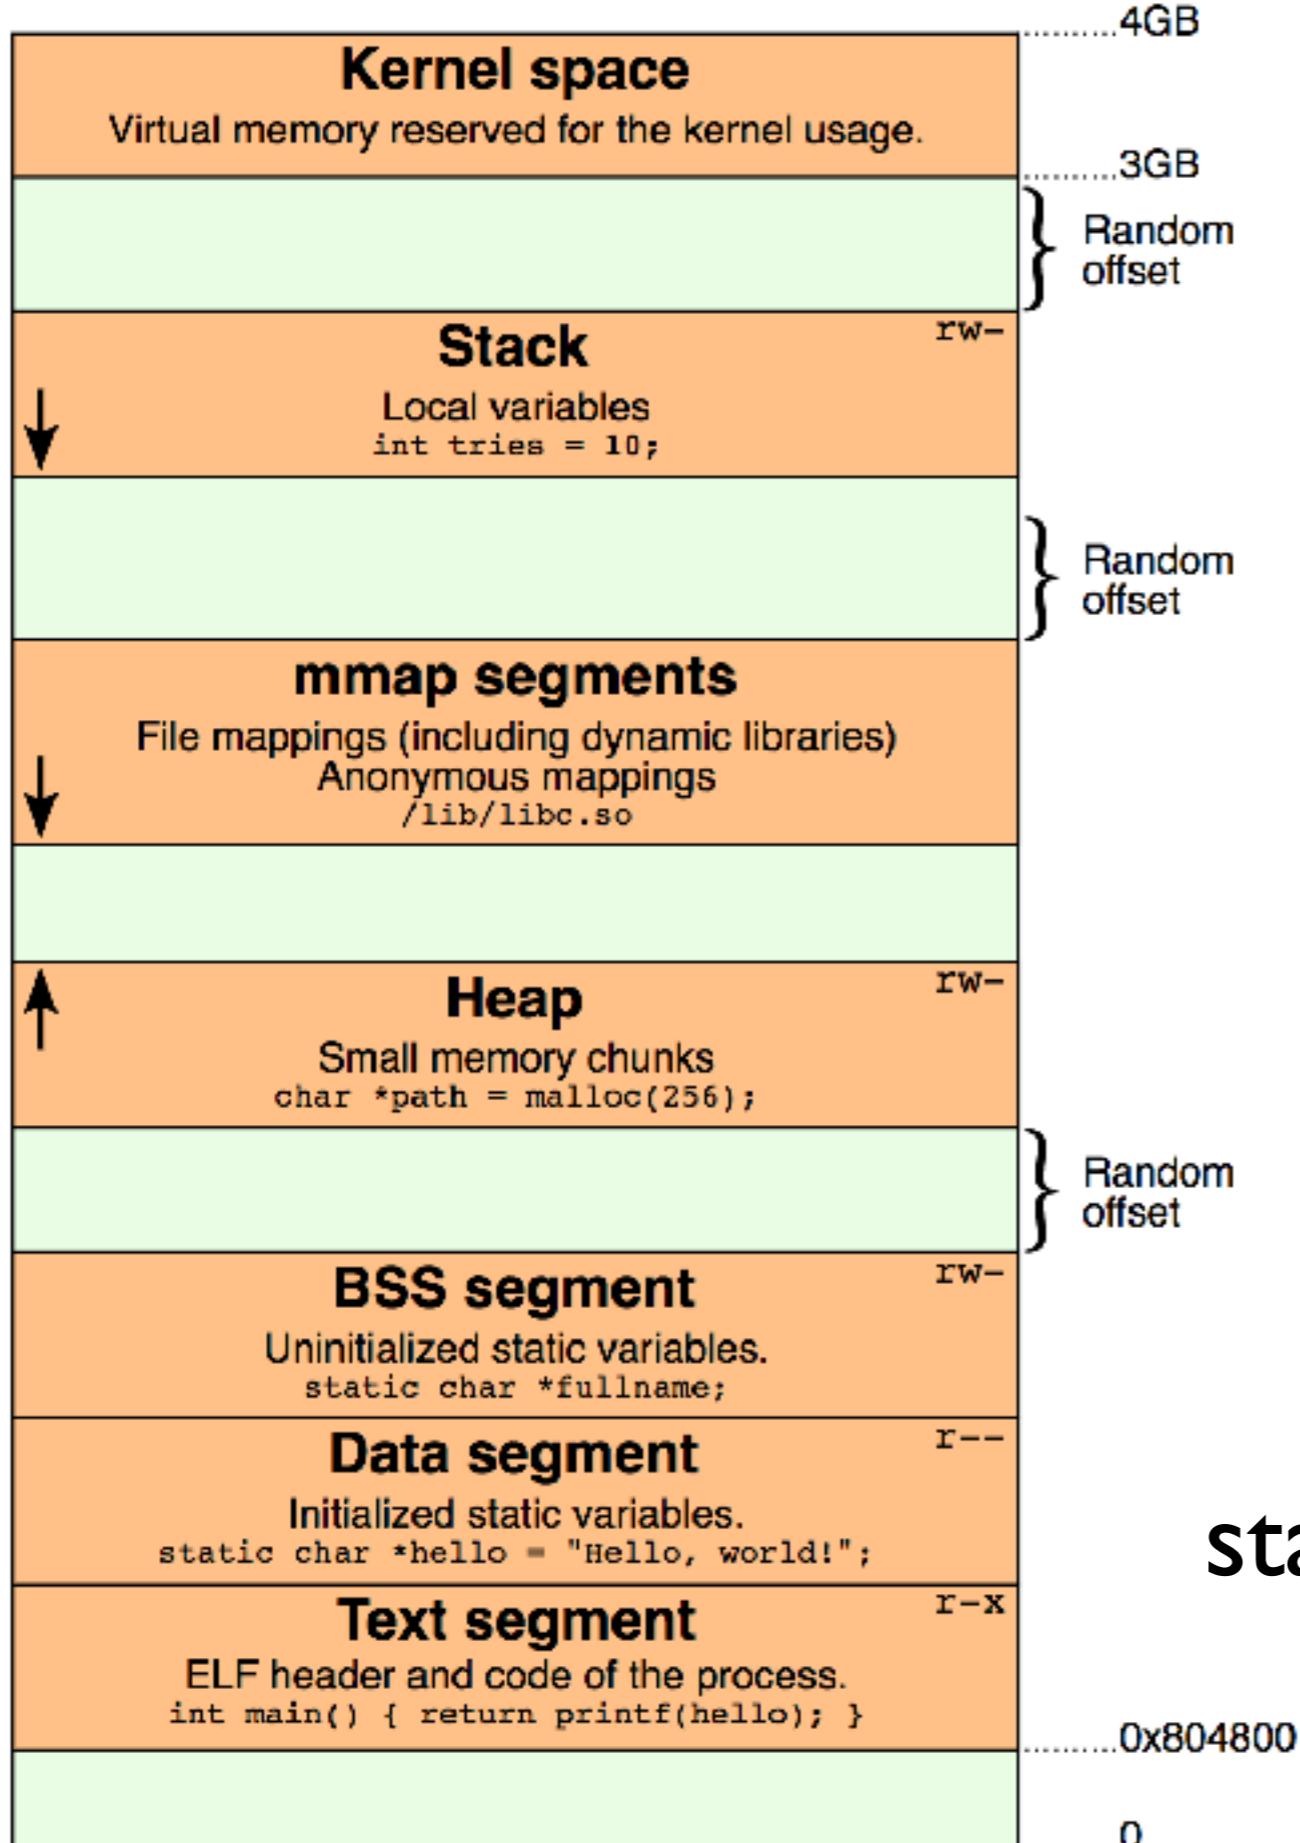
Heap: dynamic allocation

C++: New / delete

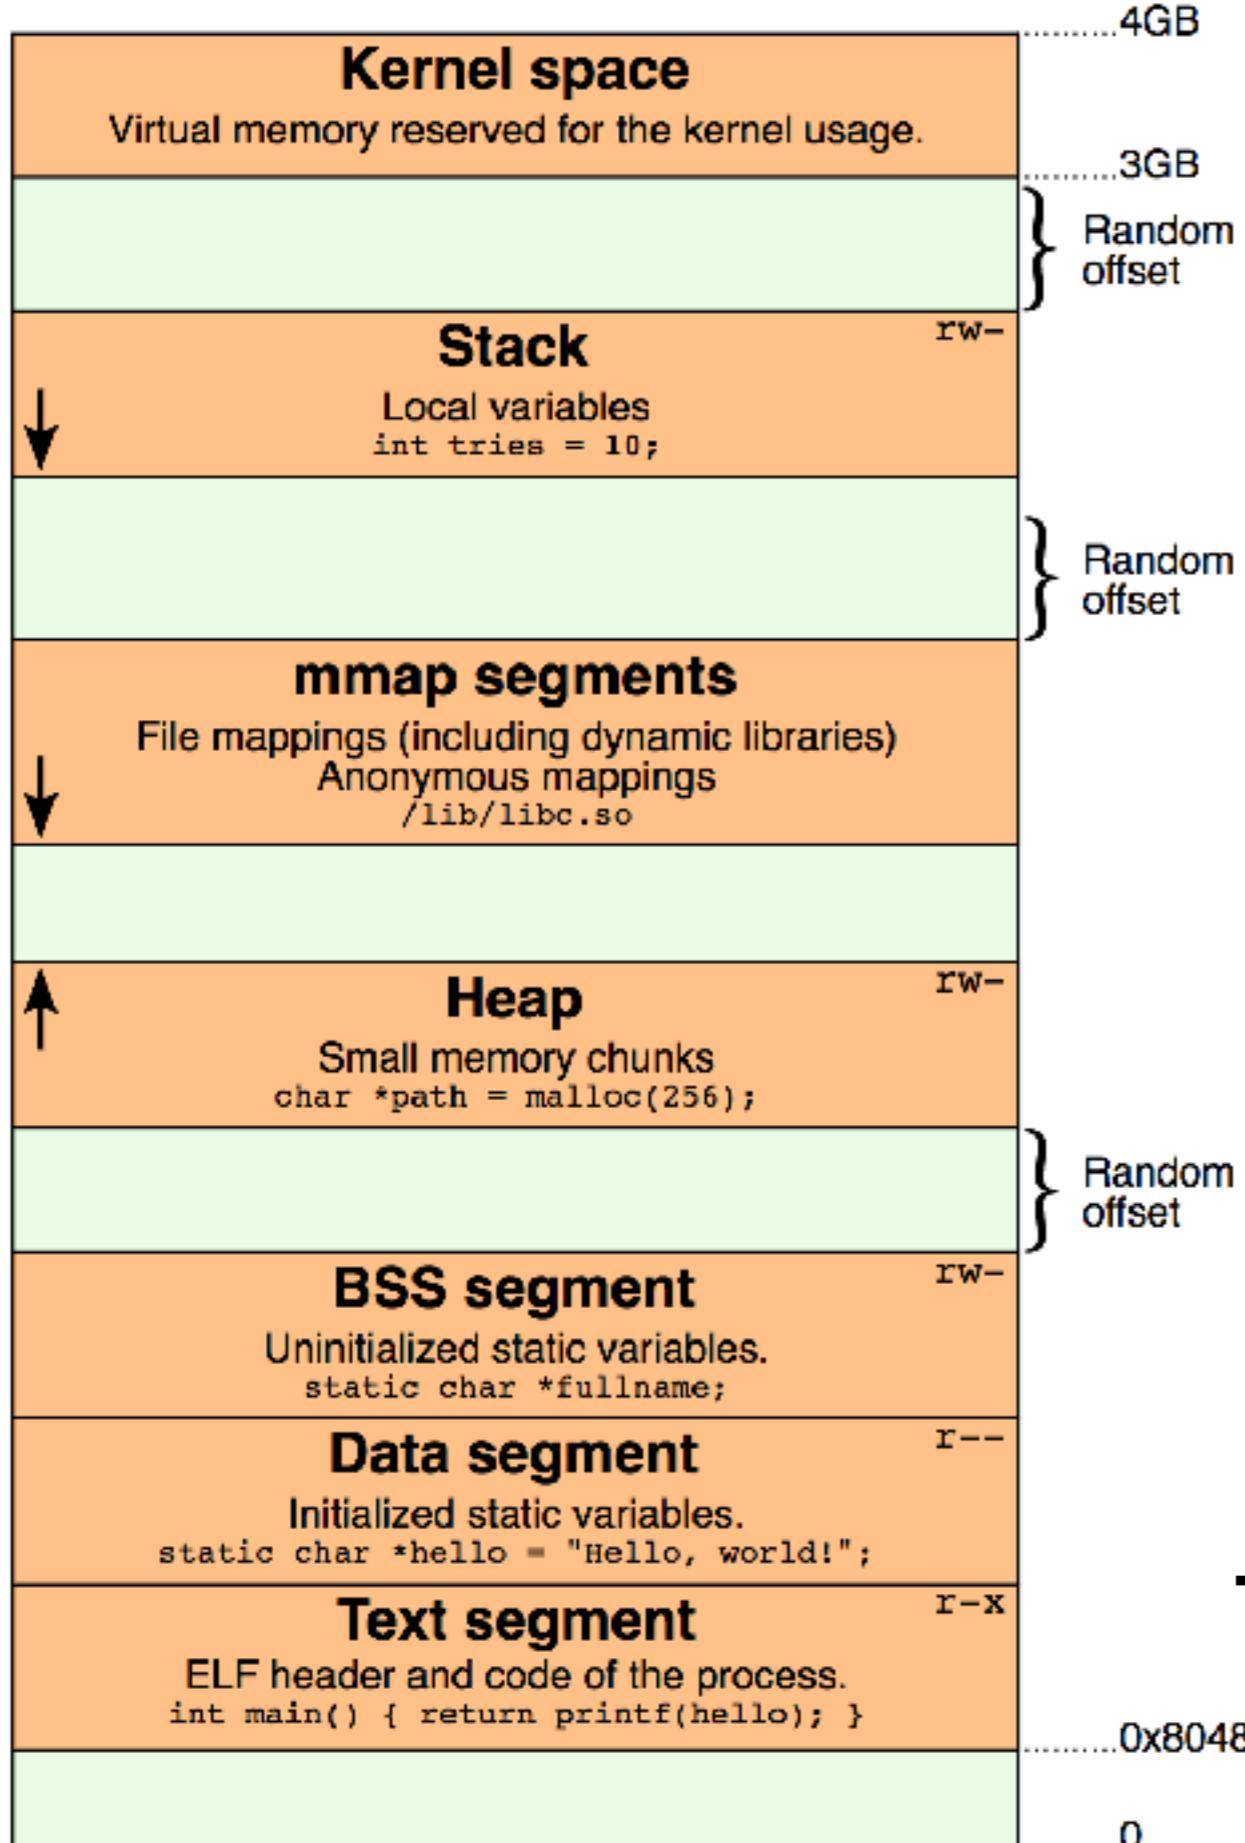
C: Malloc / free



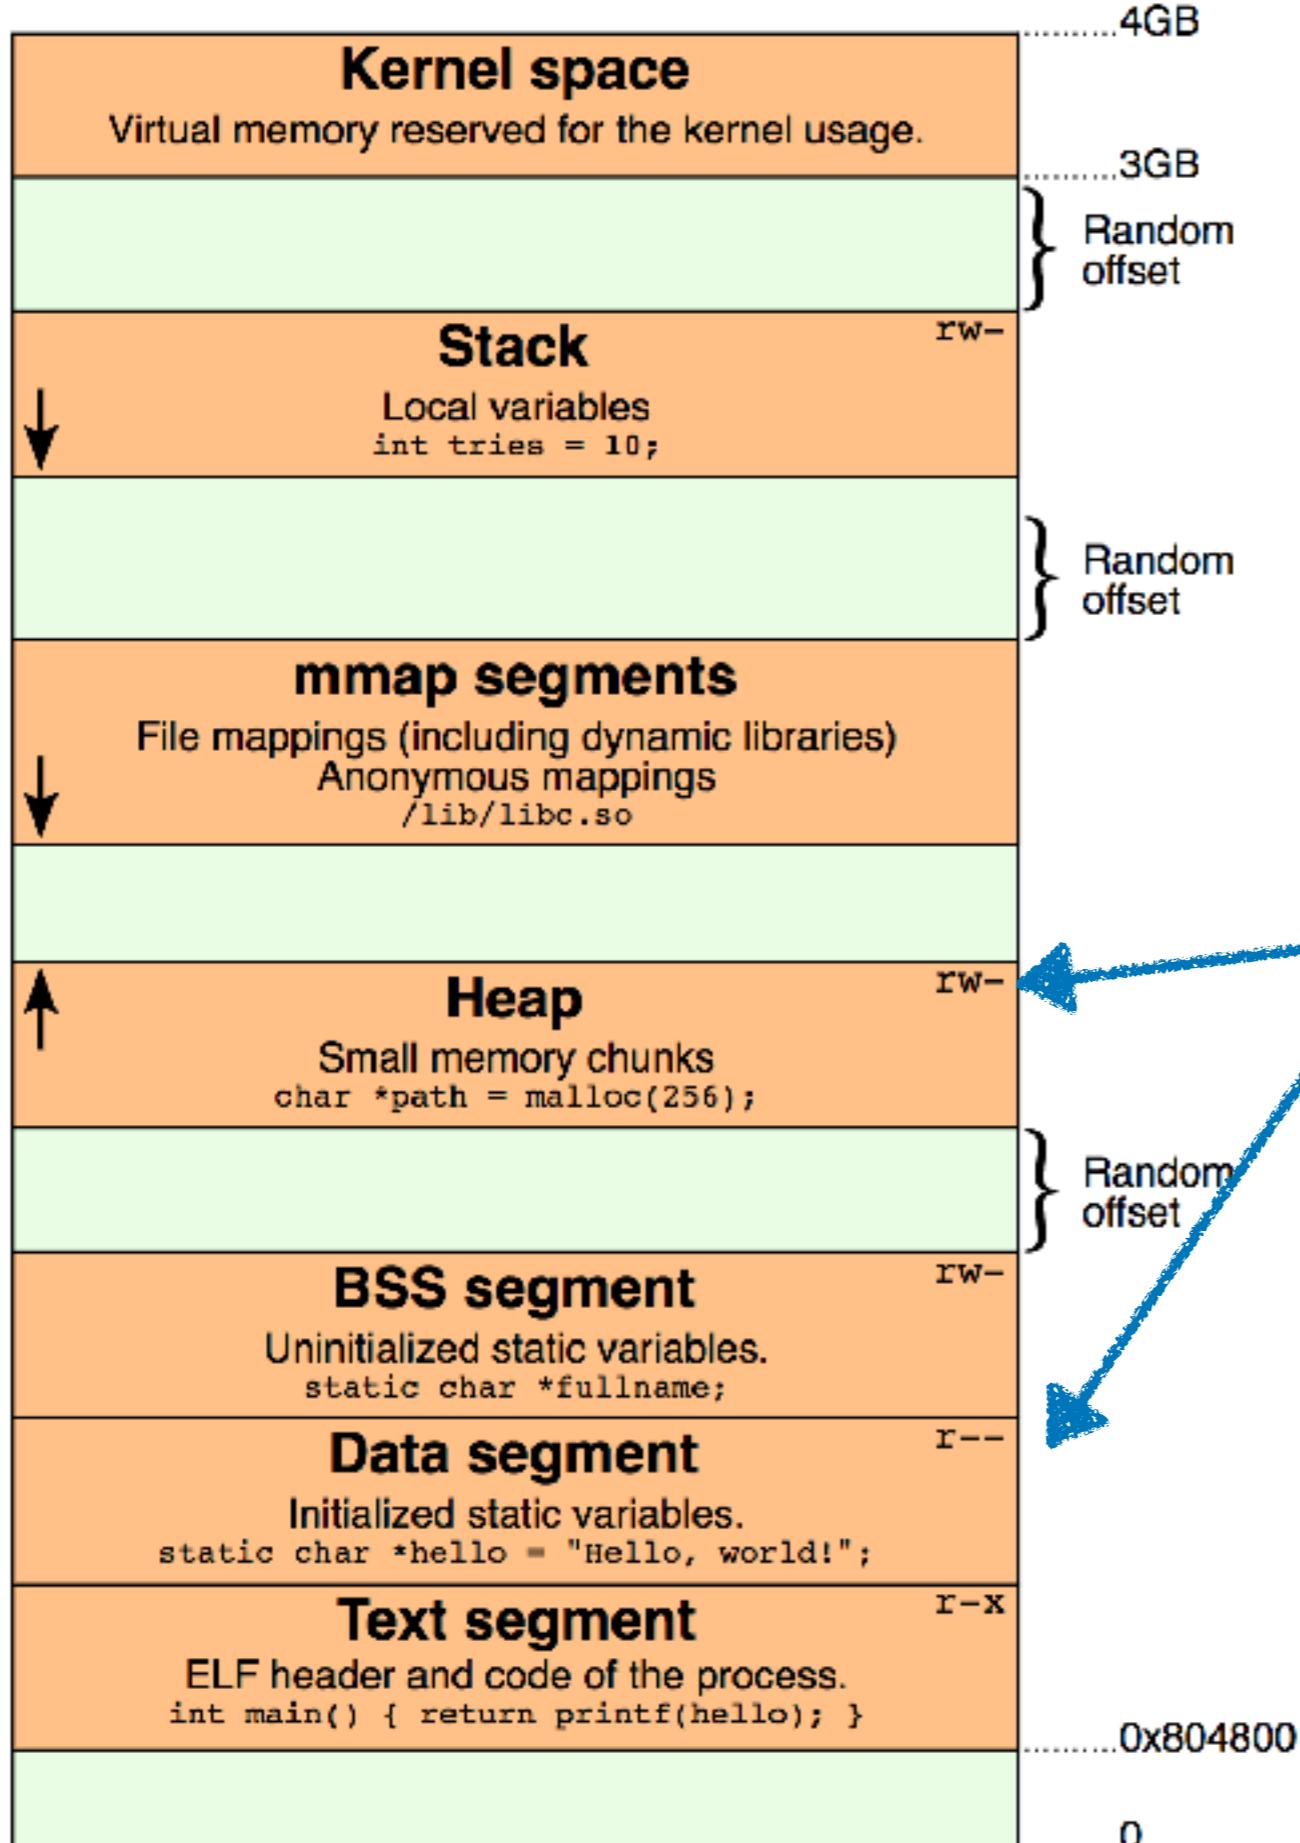
BSS: Uninitialized static vars (globals)



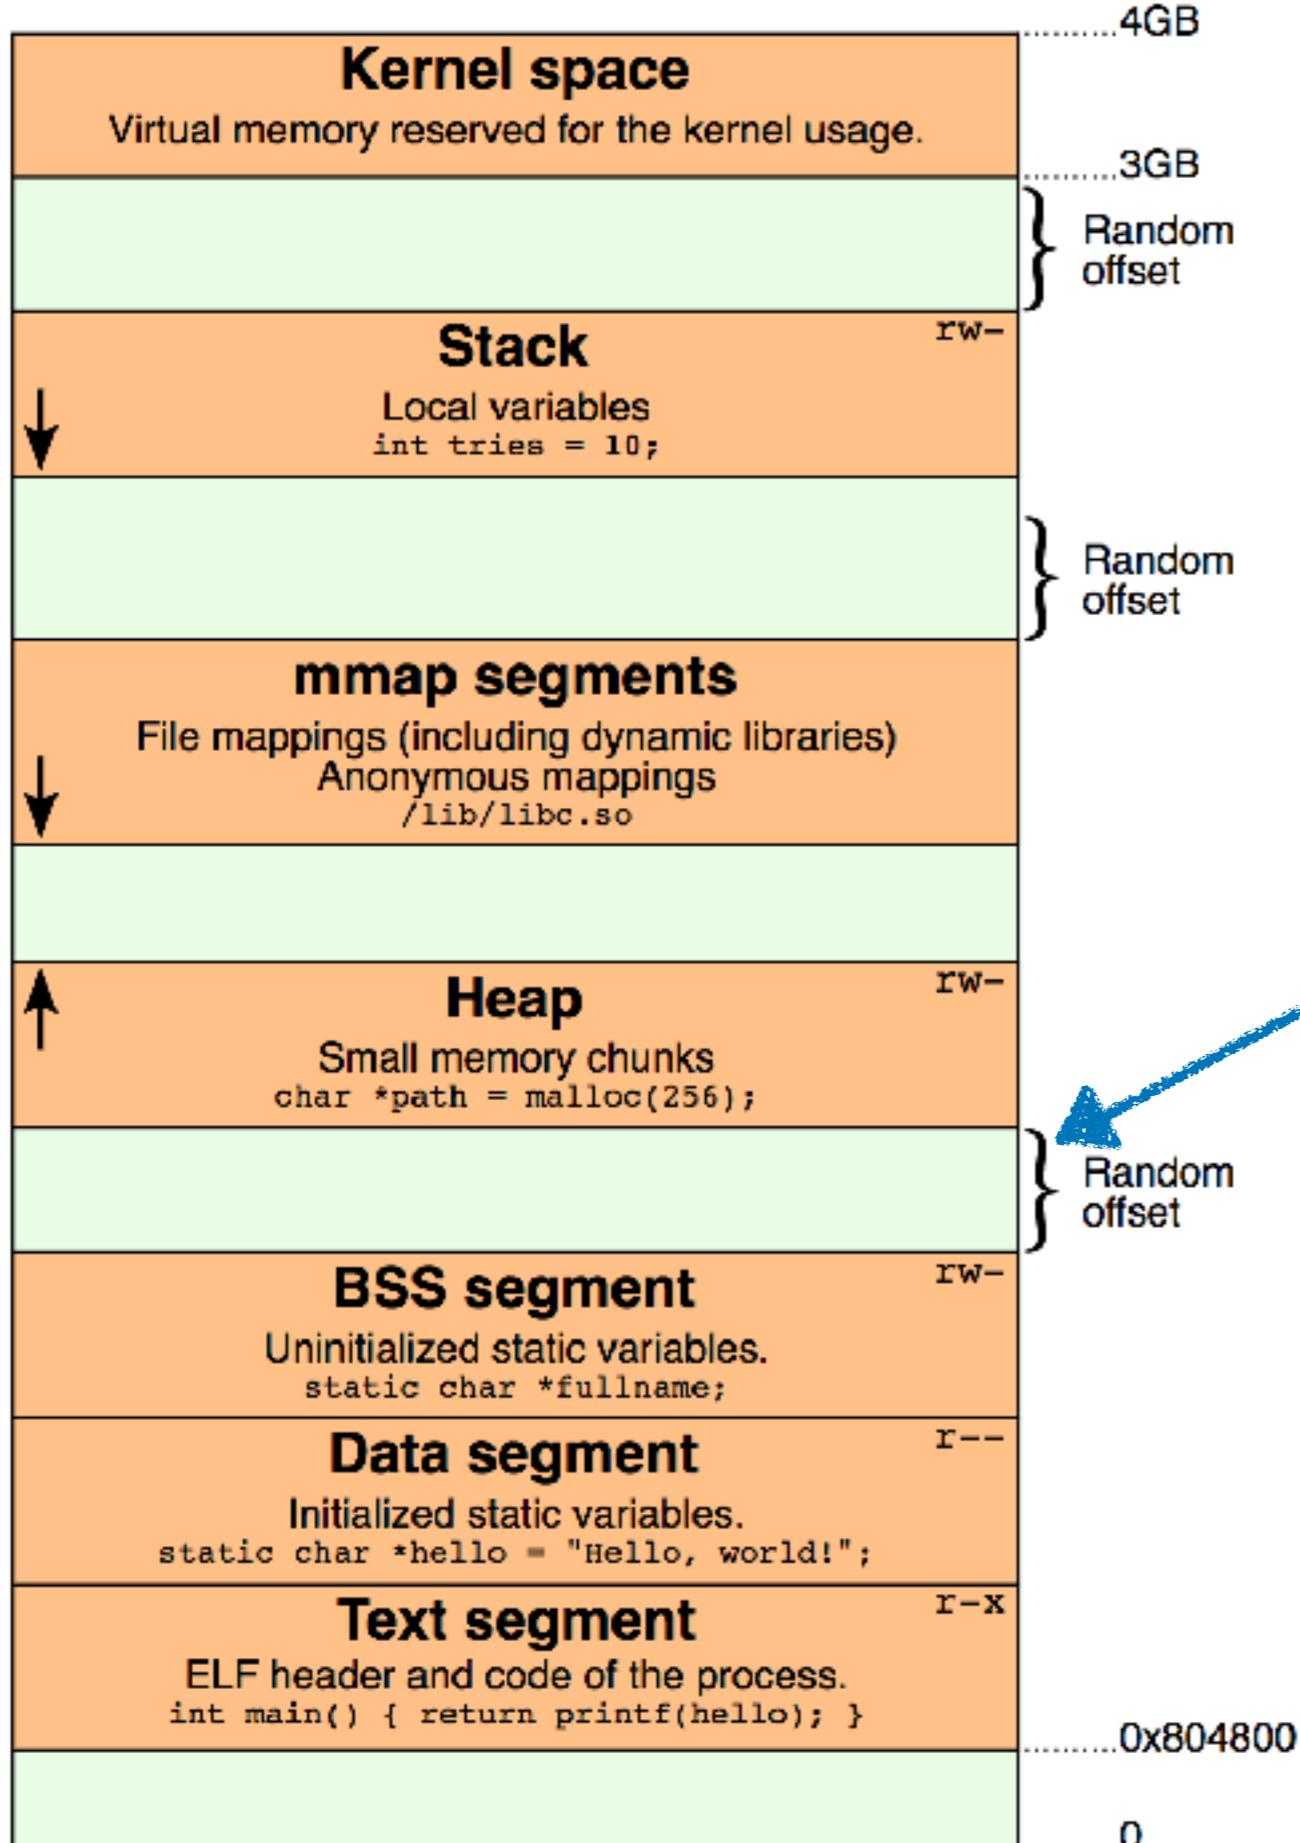
Data segment: initialized statics—e.g., constant strings



Text segment: program code



Note the permissions



This **random offset** really security feature

Calling conventions

Touch-tone phones, send an acoustic wave over the wire



If Alice wants to call Bob, her phone needs to send the right sounds over the wire in the right order

Calling conventions

When function A wants to call function B, it has to do the same

- ◆ Where do arguments go?
- ◆ How to store return address?
- ◆ Who saves registers?
- ◆ Where is result stored?

Calling conventions

Modern computers use a few **different** calling conventions

De-facto standard (Linux / MacOS / etc..) : **x86-64 System V ABI**

- Where do arguments go?
- How to store return address?
- Who saves registers?
- Where is result stored?

Note: this is **new** for the 64 bit ABI. You might see stuff online for the 32-bit ABI that is **different**

Calling conventions: x86-64

System V ABI

- ◆ Where do arguments go?
 - ◆ First six: rdi, rsi, rdx, rcx, r8, r9
- ◆ How to store return address?
 - ◆ call instruction puts on top of stack
- ◆ Who saves registers?
 - ◆ Caller saves caller-save registers
 - ◆ R10, R11, any ones used for args
- ◆ Where is result stored?
 - ◆ Result stored in %rax

x86-64 Integer Registers: Usage Conventions

%rax	Return value	%r8	Argument #5
%rbx	Callee saved	%r9	Argument #6
%rcx	Argument #4	%r10	Caller saved
%rdx	Argument #3	%r11	Caller Saved
%rsi	Argument #2	%r12	Callee saved
%rdi	Argument #1	%r13	Callee saved
%rsp	Stack pointer	%r14	Callee saved
%rbp	Callee saved	%r15	Callee saved

x86-64 System V ABI

Rules for **caller**:

- Save caller-save registers
- First six args in registers, after that put on stack
- Execute `call`—pushes ret addr

Afterwards:

- Pop saved registers
- Result now in `%rax`

x86-64 System V ABI

Rules for **callee**:

- First six args available in registers
- Push %rbp—caller's base pointer
- Move %rsp to %rbp—Setup new frame
- Subtract necessary stack space
- Push callee-save registers
- Before exit: restore rbp/callee-saved regs
 - leave instruction restores rbp
- When function done, put result in %rax
- Use ret instruction to pop return rip

These rules are cumbersome: I frequently look them up, they change depending on the kind of function you're calling, etc...

Upshot: don't feel you have to memorize, just get the gist / know how to recognize them

Small examples: interactive demo of x86-64 ABI

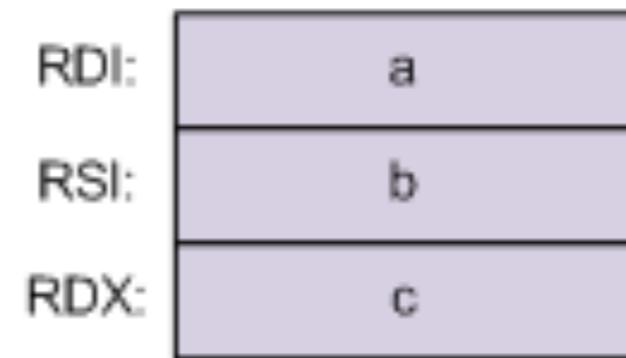
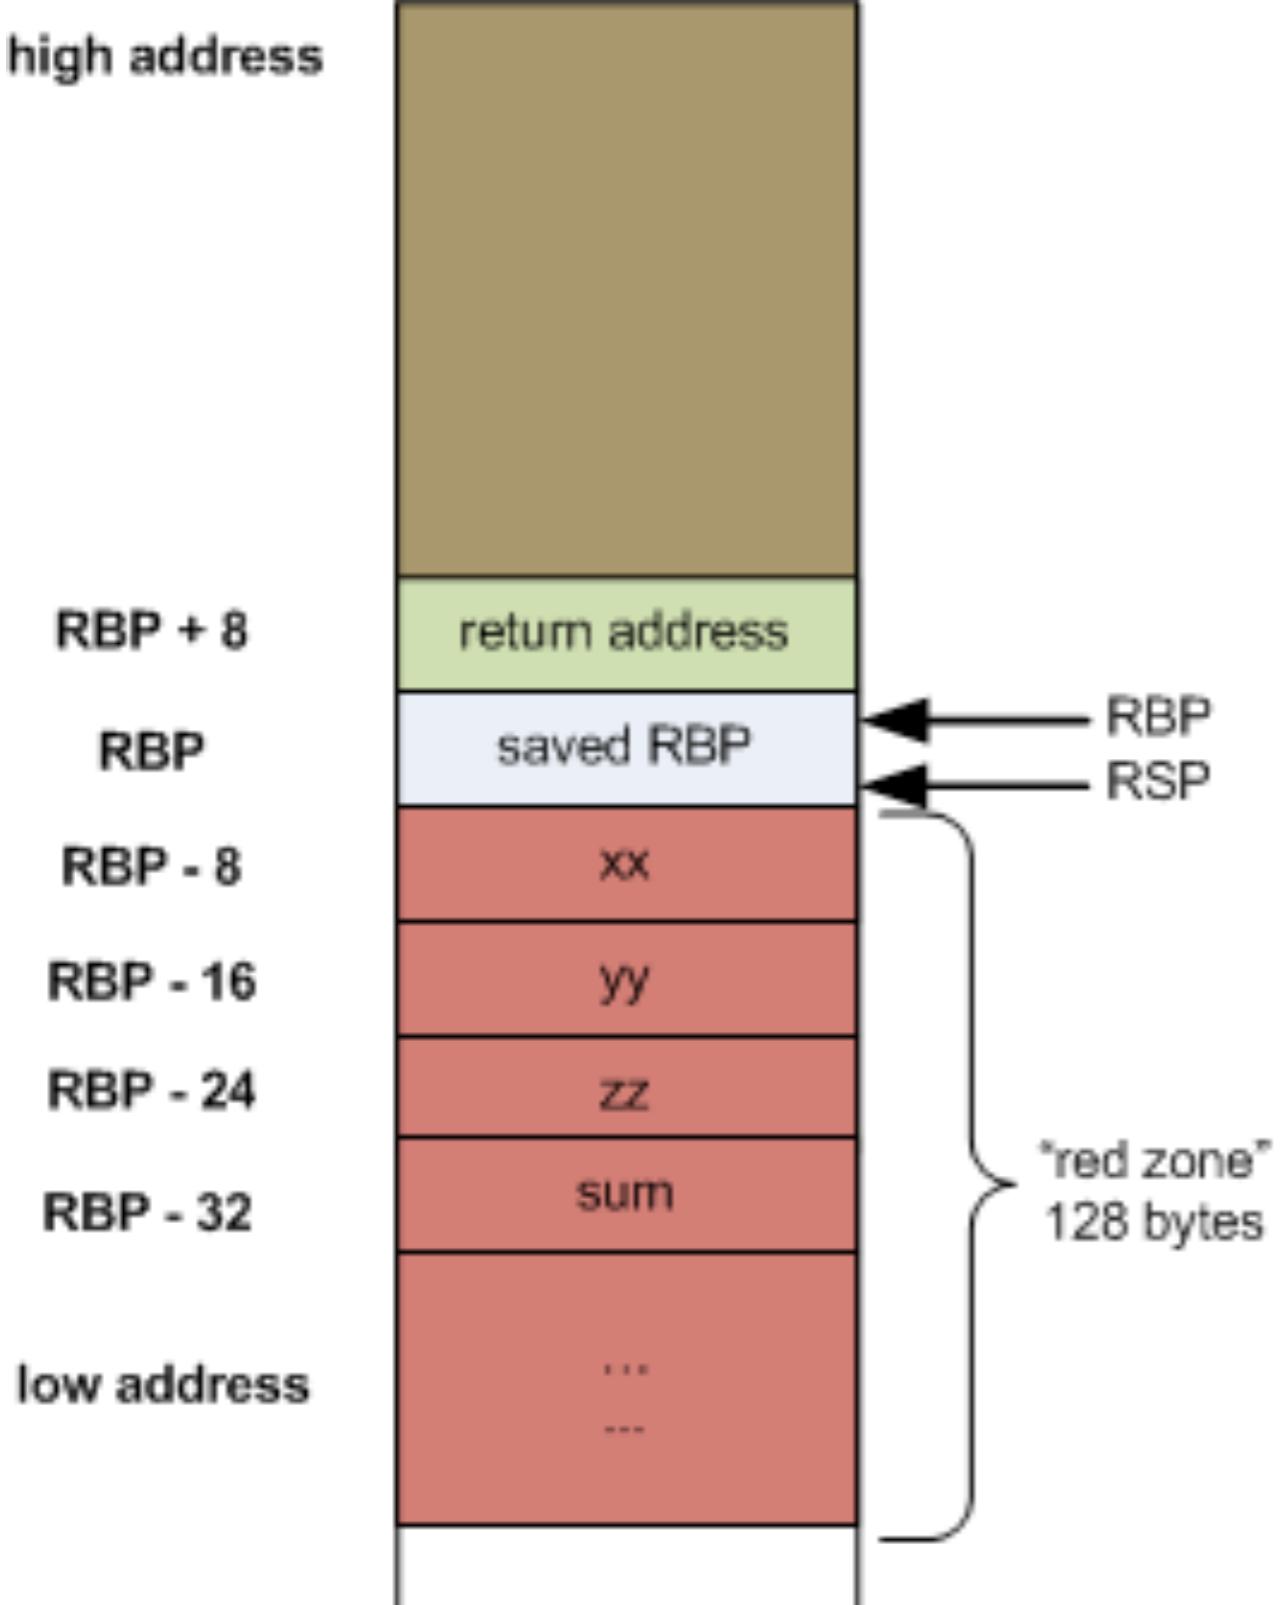
Trivia: the red zone

```
int bar(int a, int b) {  
    return a + b;  
}
```

Weird! This code uses -4(%rbp) before decrementing the stack pointer!!

Turns out: x86-64 **guarantees** there are always 128 bytes below %rsp

	bar:	
pushq		%rbp
movq		%rsp, %rbp
movl		%edi, -4(%rbp)
movl		%esi, -8(%rbp)
movl		-4(%rbp), %edx
movl		-8(%rbp), %eax
addl		%edx, %eax
popq		%rbp
ret		



Upshot: if a function uses at most 128 bytes below RSP, doesn't have to subtract anything from RSP

This is an optimization for “small” functions: so they never have to subtract from RSP

```

int foo(int x) {
    return bar(x+2);
}

int bar(int y) {
    return y*3;
}

int main(int argc, char **argv) {
    return foo(stoi(argv[0]));
}

foo: # @foo
    push    rbp
    mov     rbp, rsp
    sub     rsp, 16
    mov     dword ptr [rbp - 4], edi
    mov     eax, dword ptr [rbp - 4]
    add     eax, 2
    mov     edi, eax
    call    bar
    add     rsp, 16
    pop    rbp
    ret

bar: # @bar
    push    rbp
    mov     rbp, rsp
    mov     dword ptr [rbp - 4], edi
    imul   eax, dword ptr [rbp - 4], 3
    pop    rbp
    ret

main: # @main
    push    rbp
    mov     rbp, rsp
    sub     rsp, 16
    mov     dword ptr [rbp - 4], 0
    mov     dword ptr [rbp - 8], edi
    mov     qword ptr [rbp - 16], rsi
    mov     rax, qword ptr [rbp - 16]
    mov     rdi, qword ptr [rax]
    mov     al, 0
    call    stoi
    mov     edi, eax
    call    foo
    add     rsp, 16
    pop    rbp
    ret

```

Note: this assembly is destination-first

```

int foo(int x) {
    return bar(x+2);
}

int bar(int y) {
    return y*3;
}

int main(int argc, char **argv) {
    return foo(stoi(argv[0]));
}

```

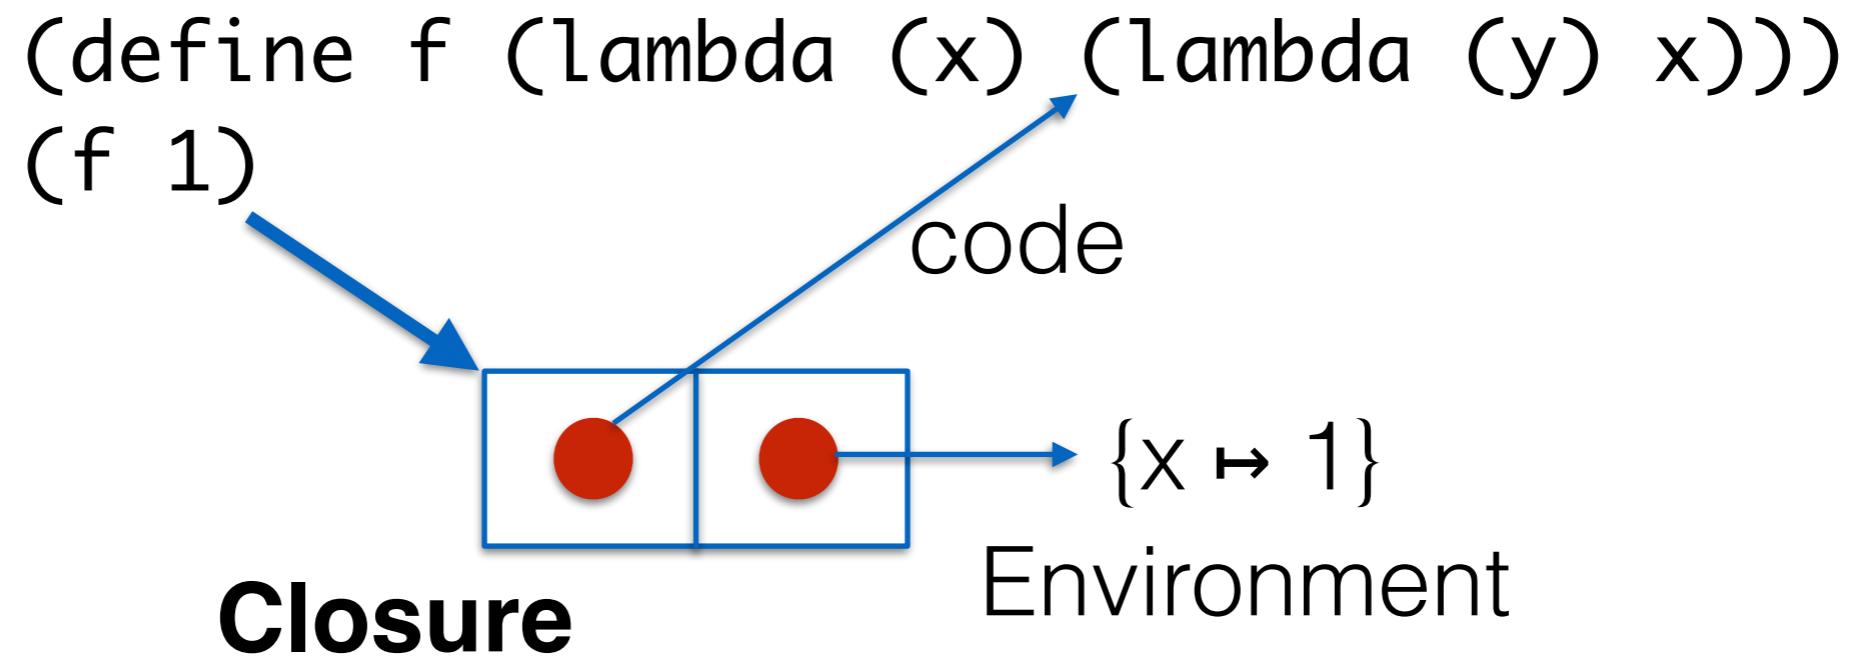
foo: add edi, 2 jmp bar	<i># @foo</i>
bar: lea eax, [rdi + 2*rdi] ret	<i># TAILCALL</i> <i># @bar</i>
main: push rax mov rdi, qword ptr [rsi] xor eax, eax call stoi mov edi, eax pop rax jmp foo	<i># @main</i>

If we turn on compiler flag -O1 (optimization-level-1) we get *much better* code

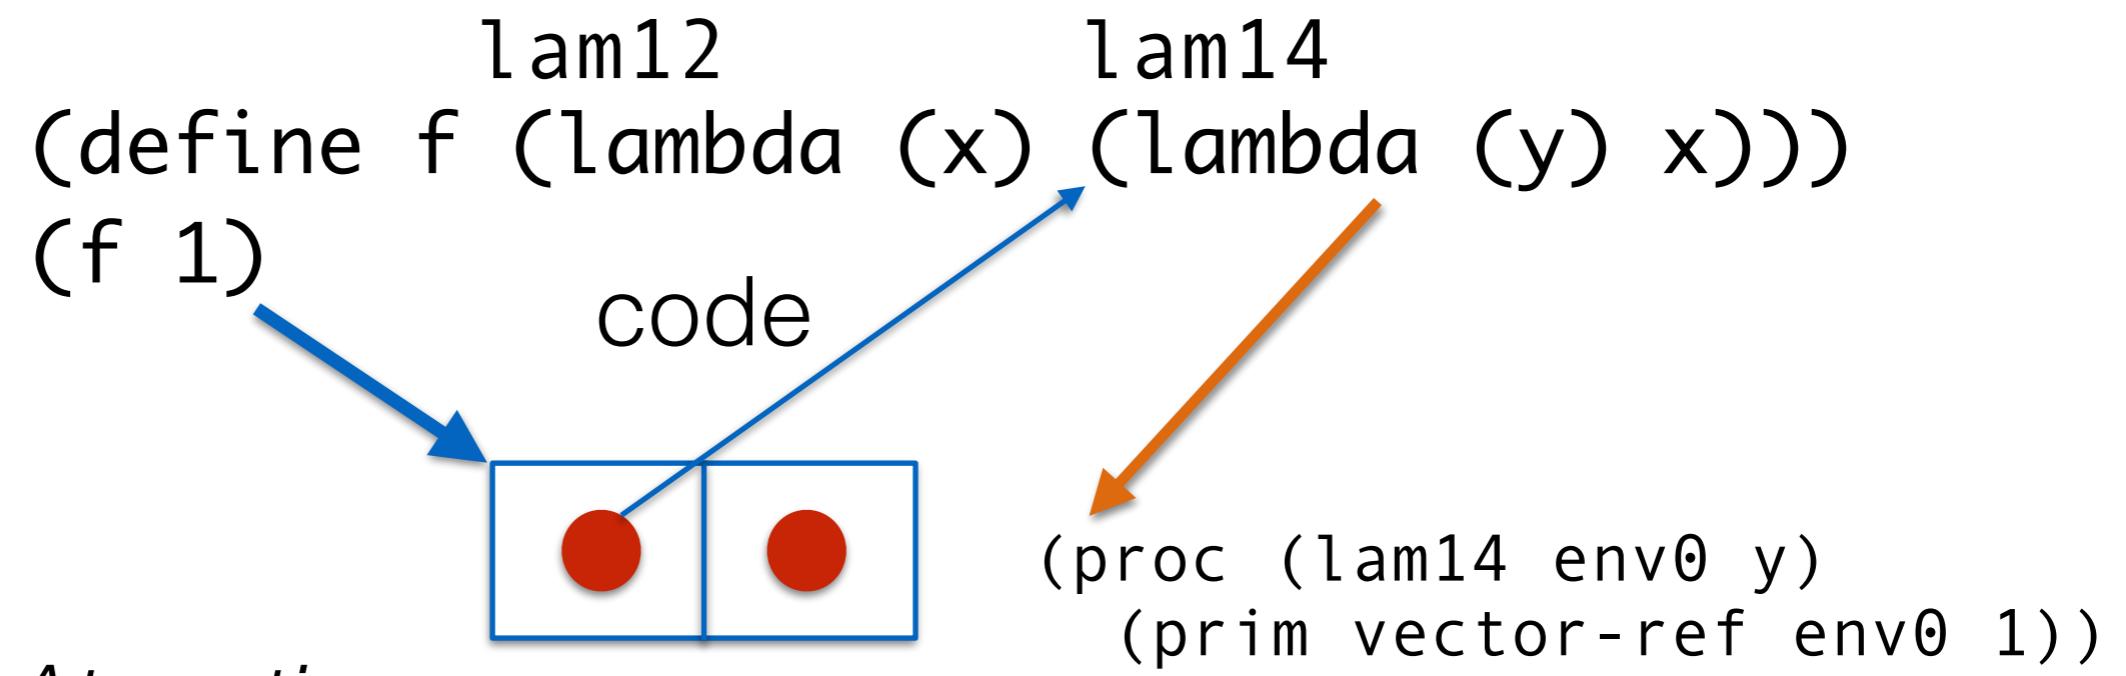
- calls to foo **and** bar are tail calls, so just turn into jumps!

**Now, back to closure
representation...**

Of course closures are code + data



In last lecture, we said we would use **vectors** to represent closures as a sequence of (a) the name of a lambda and (b) values of its free variables in a known canonical order



At runtime...

```
(vector
  lam14 ;; compiled as proc lam14
  1)    ;; remember x is offset 1
```

Then, each of these top-level procs can be compiled **directly** to C (or a C-like function)

```
(proc (lam14 env0 y)
      (prim vector-ref env0 1))
```

```
typedef unsigned long long value;
typedef value* value_vec;
typedef value_vec closure;

value lam14(closure env0, value y) {
    return env0[1];
}
```

There is a direct translation from the set of procs to C-style functions, which then translate readily to assembly...

```
lam14: # @lam14
        mov     rax, qword ptr [rdi + 8]
        ret
```

(Of course, as our IR uses CPS, all of our returns will simply be function calls)

Bottom-up closure conversion

- As the AST is traversed, free variables are computed.
- At each lambda: 1) the algorithm converts any lambdas under the lambda's body first (and also computes a set of free variables); then 2) it emits code to allocate the lambda's closure/environment and replaces free vars with env access.
- Converting the body of a lambda yields a set of free variables that can be *canonically ordered*.
- Closures are *flat* heap-allocated vectors containing a function pointer and then each free var *in order*.
- Accesses of free variables are turned into a vector-ref with the predetermined index.

...

(λ (a)

...

(λ (b c)

...

(λ (d))

...

free vars: a , c , f

(f a c)

...) ...) ...)

...

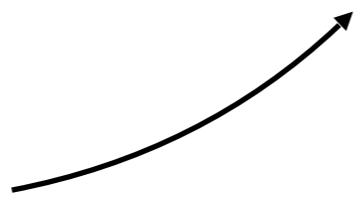
(λ (a)

...

(λ (b c)

...

adds first-order proc



allocates flat closure

...) ...) ...) ...

(proc (lam14 env d)

...

(clo-app

(prim vector-ref
env '3) ; f

(prim vector-ref
env '1) ; a

(prim vector-ref
env '2)) ; c

...)

(prim vector
lam14
a
c
f)

...

(λ (a)

...

(λ (b c)

...

(prim vector
lam14
a
c
f)

free vars: a f x y



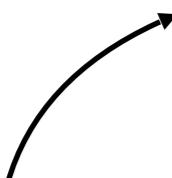
**references at closure allocation
can remain free**

...) ...) ...

```
(clo-app
  (prim vector-ref
        env '3) ;f
  (prim vector-ref
        env '1) ;a
  (prim vector-ref
        env '2)) ;c
```



```
(let ([f-clo (prim vector-ref env '3)])
  (let ([f-ptr (prim vector-ref f-clo '0)])
    (let ([a (prim vector-ref env '1)])
      (let ([c (prim vector-ref env '2)])
        (C-style-call f-ptr f-clo a c))))
```



application: 1) function pointer is accessed from closure
2) closure (f-clo) is passed to invoked function ptr

Let's look at an example that uses invocation...

Consider the following example from the lambda calculus...

```
((lambda (a) (a a))  
 (lambda (b) b))
```

Let's look at an example that uses invocation...

Consider the following example from the lambda calculus...

```
((lambda (a) (a a))  
 (lambda (b) b))
```

What would this look like post-CPS-conversion...?

Every lambda gets an extra current-continuation argument, and the (special) `halt` continuation is passed through (at runtime, `k0/k1` are `halt`)

The “final” return is within `(lambda (k1) ...)` and returns `b` to the `halt` continuation (signaling the end).

Let's look at an example that uses invocation...

Consider the following example from the lambda calculus...

```
((lambda (a) (a a))  
 (lambda (b) b))
```

What would this look like post-CPS-conversion...?

```
((lambda (k0 a) (a k0 a))  
 halt  
 (lambda (k1 b) (k1 b '0)))
```

Every lambda gets an extra current-continuation argument, and the (special) `halt` continuation is passed through (at runtime, `k0/k1` are `halt`)

The “final” return is within `(lambda (k1) ...)` and returns `b` to the `halt` continuation (signaling the end).

What will closure conversion then produce...?

```
lam5
  ((lambda (k0 a) (a k0 a))
   lam16 halt
   (lambda (k1 b) (k1 b '0)))
```

```
(proc (lam5 env0 k0 a) (proc (lam16 env1 k1 b)
  (clo-app a k0 a)) (clo-app k1 b '0))
```

Let's live code bottom-up closure conversion.

```
; Input Language (simplified cps):
; e ::= (let ([x 'dat]) e)
;       | (let ([x (prim op x ...)]) e)
;       | (let ([x (lambda (x ...) e)]) e)
;       | (if x e e)
;       | (x x ...)
```

Our output language is a list of C-style procedures

```
; Output Language (procedural):
; p ::= ((proc (x x ...) e) ...)
; e ::= (let ([x 'dat]) e)
;       | (let ([x (prim op x ...)]) e)
;       | (let ([x (make-closure x x ...)]) e)
;       | (let ([x (env-ref x nat)]) e)
;       | (if x e e)
;       | (clo-app x x ...)
```

```
(define (closure-convert simple-cps)
  ; ; Exp x List[Proc] -> Exp x Set[Var] x List[Proc]
  (define (bottom-up e procs)
    (match e
      ...))
  ; ; returns the top expression, set of free variables
  ; ; and list of sub-procs
  (match-define `(,main-body ,free ,procs)
    (bottom-up simple-cps '()))
  ; ; return the list of procs & main expression proc
  `((proc (main) ,main-body) . ,procs))
```

```
(define (closure-convert simple-cps)
;; Exp x List[Proc] -> Exp x Set[Var] x List[Proc]
(define (bottom-up e procs)
  (match e
    [ `(let ([ ,x ' ,dat]) ,e0)
     (match-define `( ,e0+ ,freet+ ,procs+)
                  (bottom-up e0 procs))
     `((let ([ ,x ' ,dat]) ,e0+)
       ,(set-remove free+ x)
       ,procs+)]
```

```
(define (closure-convert simple-cps)
  ;; Exp x List[Proc] -> Exp x Set[Var] x List[Proc]
(define (bottom-up e procs)
  (match e
    ...
    [ `(let ([,x (prim ,op ,xs ...)]) ,e0)
      (match-define `(,e0+ ,free+ ,procs+)
                    (bottom-up e0 procs))
      `((let ([,x (prim ,op ,@xs)]) ,e0+)
        ,(set-remove (set-union free+ (list->set xs)) x)
        ,procs+)]
```

```

(define (closure-convert simple-cps)
;; Exp x List[Proc] -> Exp x Set[Var] x List[Proc]
(define (bottom-up e procs)
  (match e
    ...
    [`(let ([,x (lambda (,xs ...) ,body)]) ,e0)
     (match-define `(,e0+ ,free0+ ,procs0+)
                  (bottom-up e0 procs))
     (match-define `(,body+ ,freelam+ ,procs1+)
                  (bottom-up body procs0+))
     (define env-vars (foldl (lambda (x fr) (set-remove fr x))
                             freelam+
                             xs))
     (define ordered-env-vars (set->list env-vars))
     (define lamx (gensym 'lam))
     (define envx (gensym 'env))
     (define body++ (cdr (foldl (lambda (x count+body)
                                   (match-define (cons cnt bdy) count+body)
                                   (cons (+ 1 cnt)
                                         `(let ([,x (env-ref ,envx ,cnt)])
                                            ,bdy)))
                                   (cons 1 body+)
                                   ordered-env-vars))))
     `( (let ([,x (make-closure ,lamx ,@ordered-env-vars)]) ,e0+)
        ,(set-remove (set-union free0+ env-vars) x)
        ((proc (,lamx ,envx ,@xs) ,body++) . ,procs1+))]
```

```
(define (closure-convert simple-cps)
;; Exp x List[Proc] -> Exp x Set[Var] x List[Proc]
(define (bottom-up e procs)
  (match e
    ...
    [ `(if ,(? symbol? x) ,e0 ,e1)
      (match-define `(,e0+ ,free0+ ,procs0+)
                    (bottom-up e0 procs))
      (match-define `(,e1+ ,free1+ ,procs1+)
                    (bottom-up e1 procs0+))
      `((if ,x ,e0+ ,e1+)
        ,(set-union free1+ free0+ (set x))
        ,procs1+)]
```

```
(define (closure-convert simple-cps)
;; Exp x List[Proc] -> Exp x Set[Var] x List[Proc]
(define (bottom-up e procs)
  (match e
    ...
    [ `( , (? symbol? xs) . . . )
      `((clo-app ,@xs)
        ,(list->set xs)
        ,procs) ] )
```